

Week 3

Synchronization Primitives

Max Kopinsky
21 January, 2025

Teams Registration is Open

Content Zoom Lecture Recordings Discussions **Assignments** Classlist Grades More ▾

Assignments

[Help](#)

New Assignment

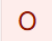

Edit Categories

More Actions ▾

[Bulk Edit](#)

<input type="checkbox"/>	Assignment	New Submissions	Completed	Evaluated	Feedback Published
	No Category				
<input type="checkbox"/>	Teams Registration ▾ Due on Jan 28, 2025 11:59 PM	3	3/370	0/370	0/370

Max Kopinsky / operating-systems-w25

 **operating-systems-w25** 

 ▾

 Star

0

 Fork

3

⋮

Project information

 2 Commits  1 Branch  0 Tags  3 KiB Project Storage


 README  Auto DevOps enabled [+ Add LICENSE](#) [+ Add CHANGELOG](#) [+ Add CONTRIBUTING](#)

[+ Add Kubernetes cluster](#) [+ Add Wiki](#) [+ Configure Integrations](#)

Created on

January 20, 2025

 main ▾


operating-systems-w25 /  ▾

History

Find file

Edit ▾

Code ▾

 **initial commit**

Max Kopinsky authored 2 hours ago

27858953



Name

Last commit

Last update

 README.md

initial commit

2 hours ago

Teams Registration is Open

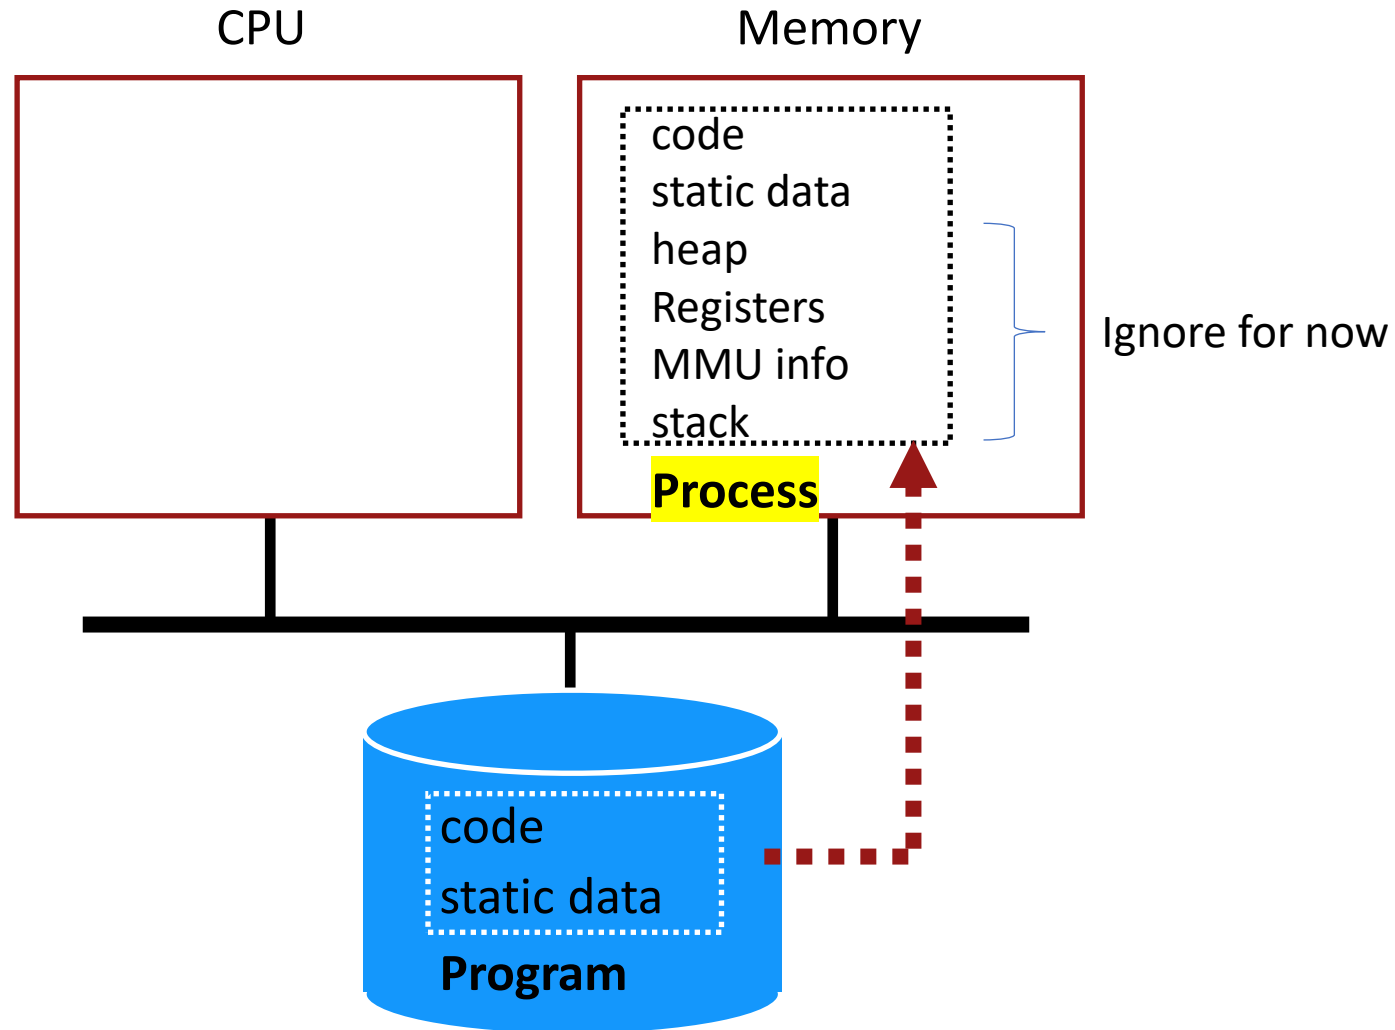
- You **must** fork the repository. You won't be able to submit your work to my repository. You can change the name of your fork if you want.
- If you want to work alone, you must submit a team registration with no teammate. Otherwise our autograder doesn't know about you.
- Your registration must be a **plain text** file with some information about your team. **Read the directions carefully.** The submissions are processed by a script. You will consume **a lot** of TA time if we have to track you down to correct errors.

Recap from Week 2

- Process
- Linux process tree
- Process switch
- Process scheduler

Recap from Week 2

Process = Program in execution




Recap from Week 2

Linux Process Primitives

- `pid = fork()`
- `exec(filename)`
- `exit()`
- `wait()`


Recap from Week 2

Process = Environment + Code




Recap from Week 2

After a fork()



Recap from Week 2

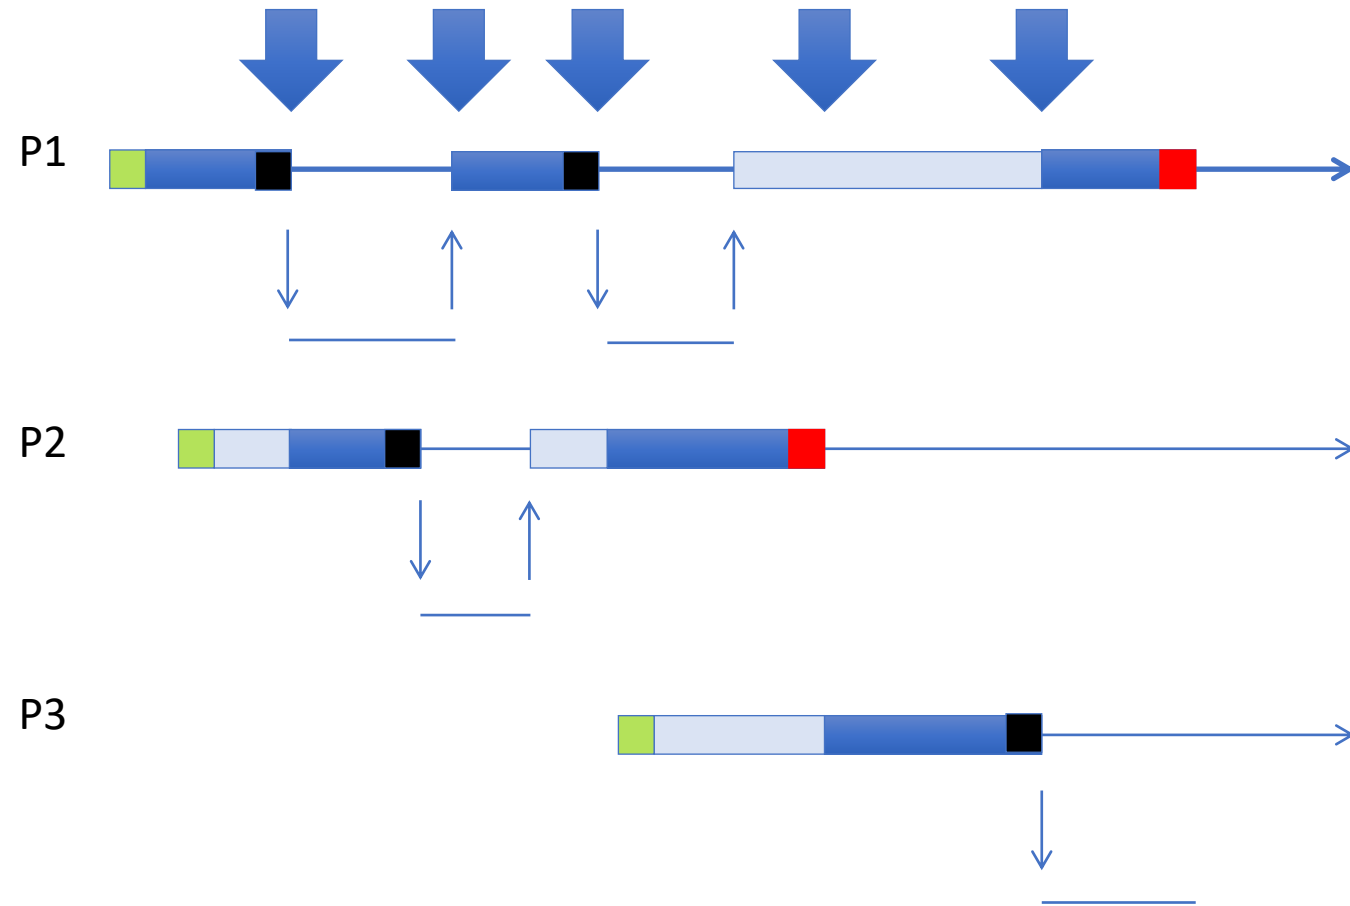
After an exec() in the Child



Recap from Week 2

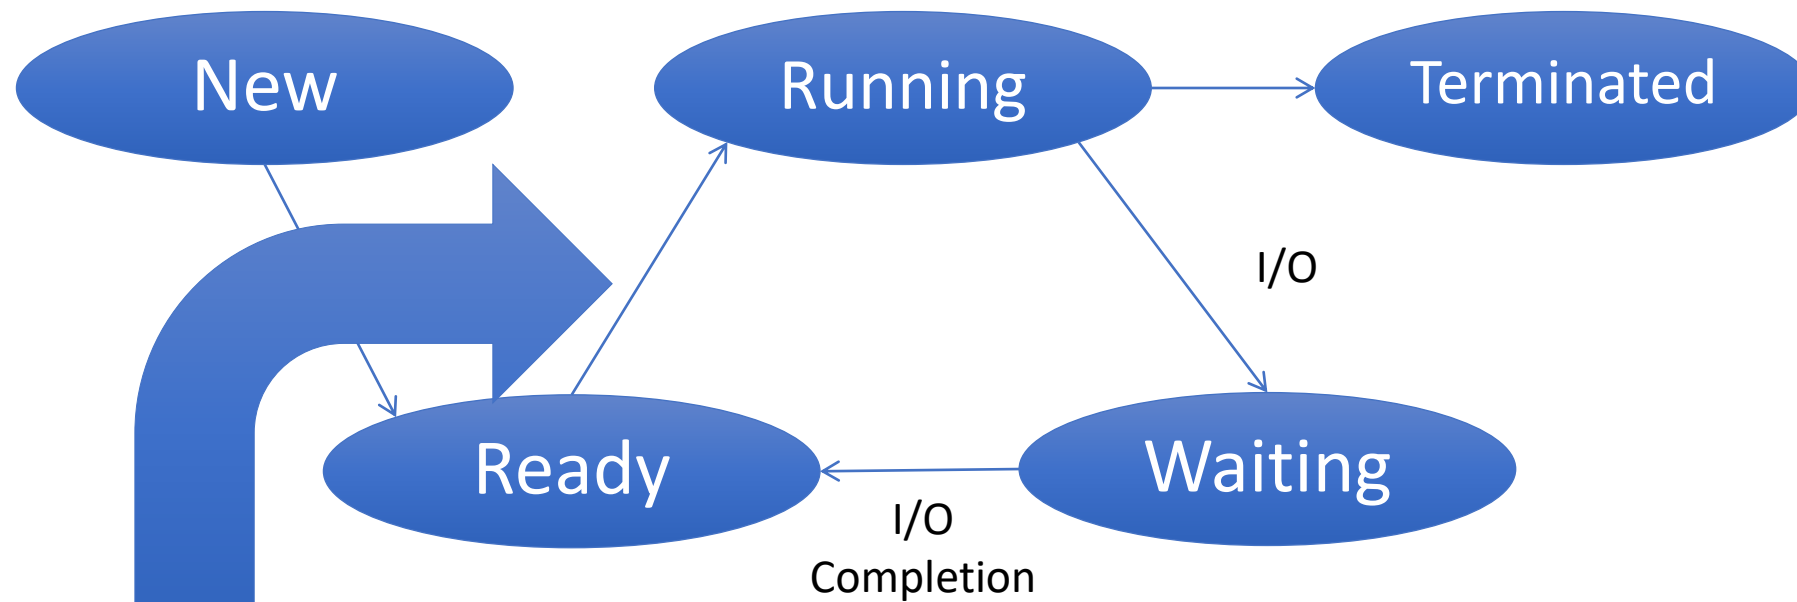
Process Switch

- Change of process using the CPU
- Save and restore registers and other info



Recap from Week 2

Process Scheduling



Many processes may be ready.
Process scheduler picks one.

Questions from last week?

Before we begin with today's topic

- Concurrency is a large sub-field of computer science
- In this course, we get a small taste of it
- If you enjoy this lecture, consider [COMP-409 Concurrent programming](#)
 - Highly recommend for a strong systems background

Real world concurrency

- Millions of drivers on highway at once.
- Student does homework while watching Netflix.
- Faculty has lunch while grading papers


Real world concurrency

- Millions of drivers on highway at once.
- Student does homework while watching Netflix.
- Faculty has lunch while grading papers... and watching Netflix.

Key Concepts for Today


- Process vs Thread
- Mutual Exclusion
- Locking
- Deadlocks
- Conditional variables

Motivation for Concurrency – Moore's Law




Motivation for Concurrency

- CPU trend: Same speed, but multiple cores
- Goal: write applications that fully utilize many cores




Concurrency Abstraction vs Reality


Programmer Abstraction
(limitless virtualized CPUs)



Physical Reality
(small # of physical CPUs)



Ready queue




Concurrency – Option 1

- Build apps from many communicating **processes**
- Communicate through **message passing**
 - No shared memory
- Pros
 - If one process crashes, other processes unaffected
- Cons
 - High communication overheads
 - Expensive context switching

Concurrency – Option 1

- Build apps from many communicating **processes**
- Communicate through **message passing**
 - No shared memory
- Pros
 - If one process crashes, other processes unaffected
- Cons
 - High communication overheads
 - Expensive context switching



Will see
next week

Concurrency – Option 2

- New abstraction: **thread**
- Multiple threads in a process
- Threads are like processes except
 - Multiple threads in the same process share an address space
 - **Communicate through shared address space**
 - If one thread crashes,
 - the entire process, including other threads, crashes

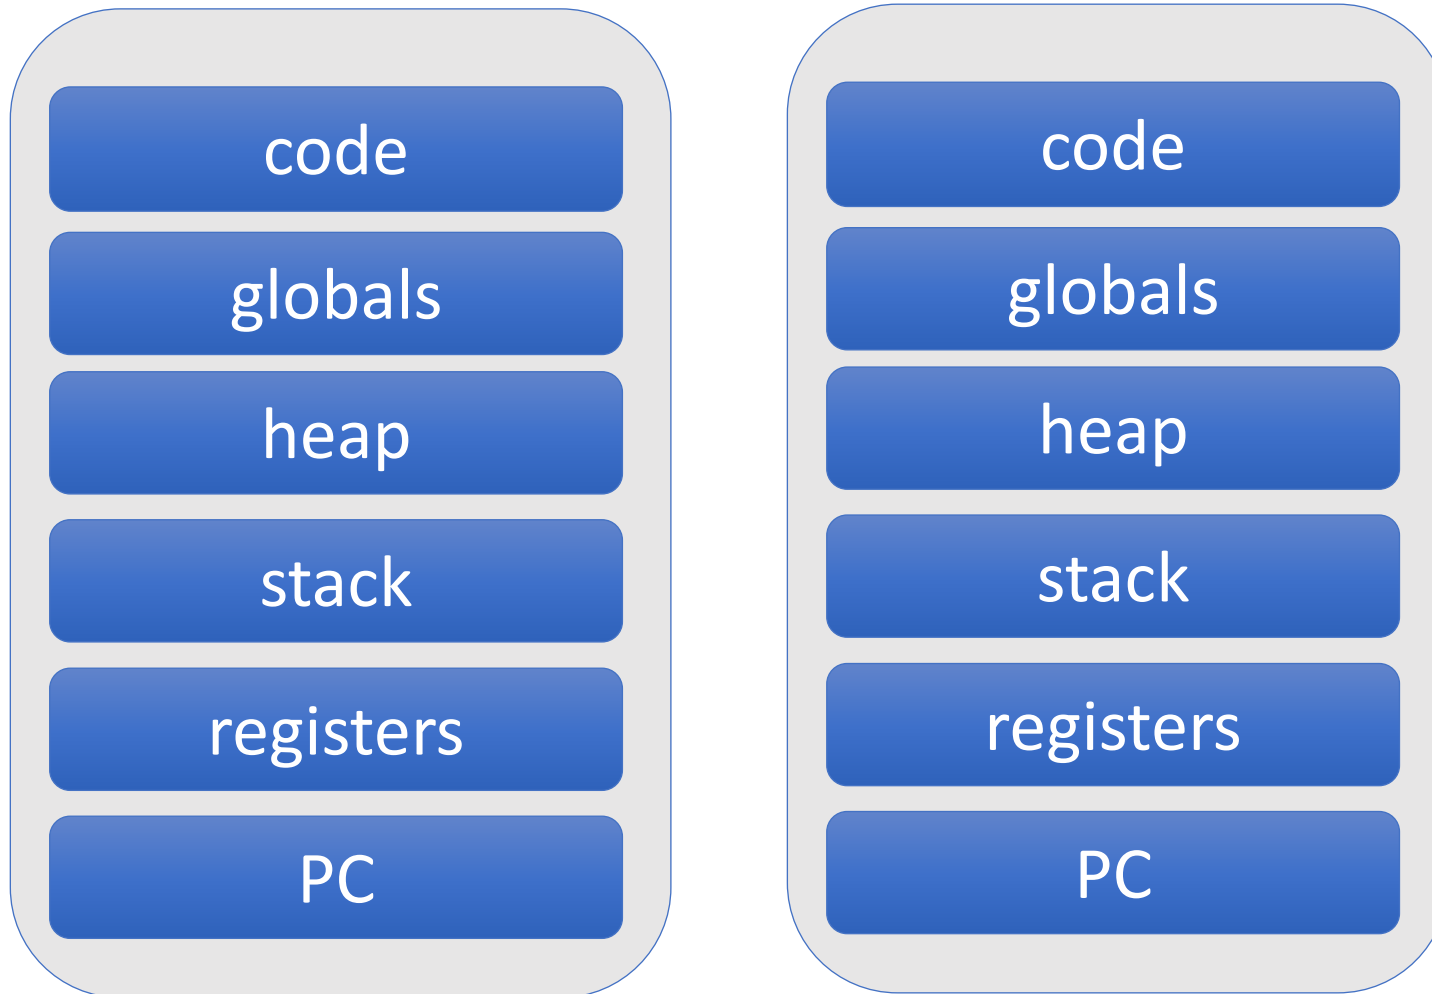
Concurrency – Option 2

- New abstraction: **thread**
- Multiple threads in a process
- Threads are like processes except
 - Multiple threads in the same process share an address space
 - **Communicate through shared address space**
 - If one thread crashes,
 - the entire process, including other threads, crashes

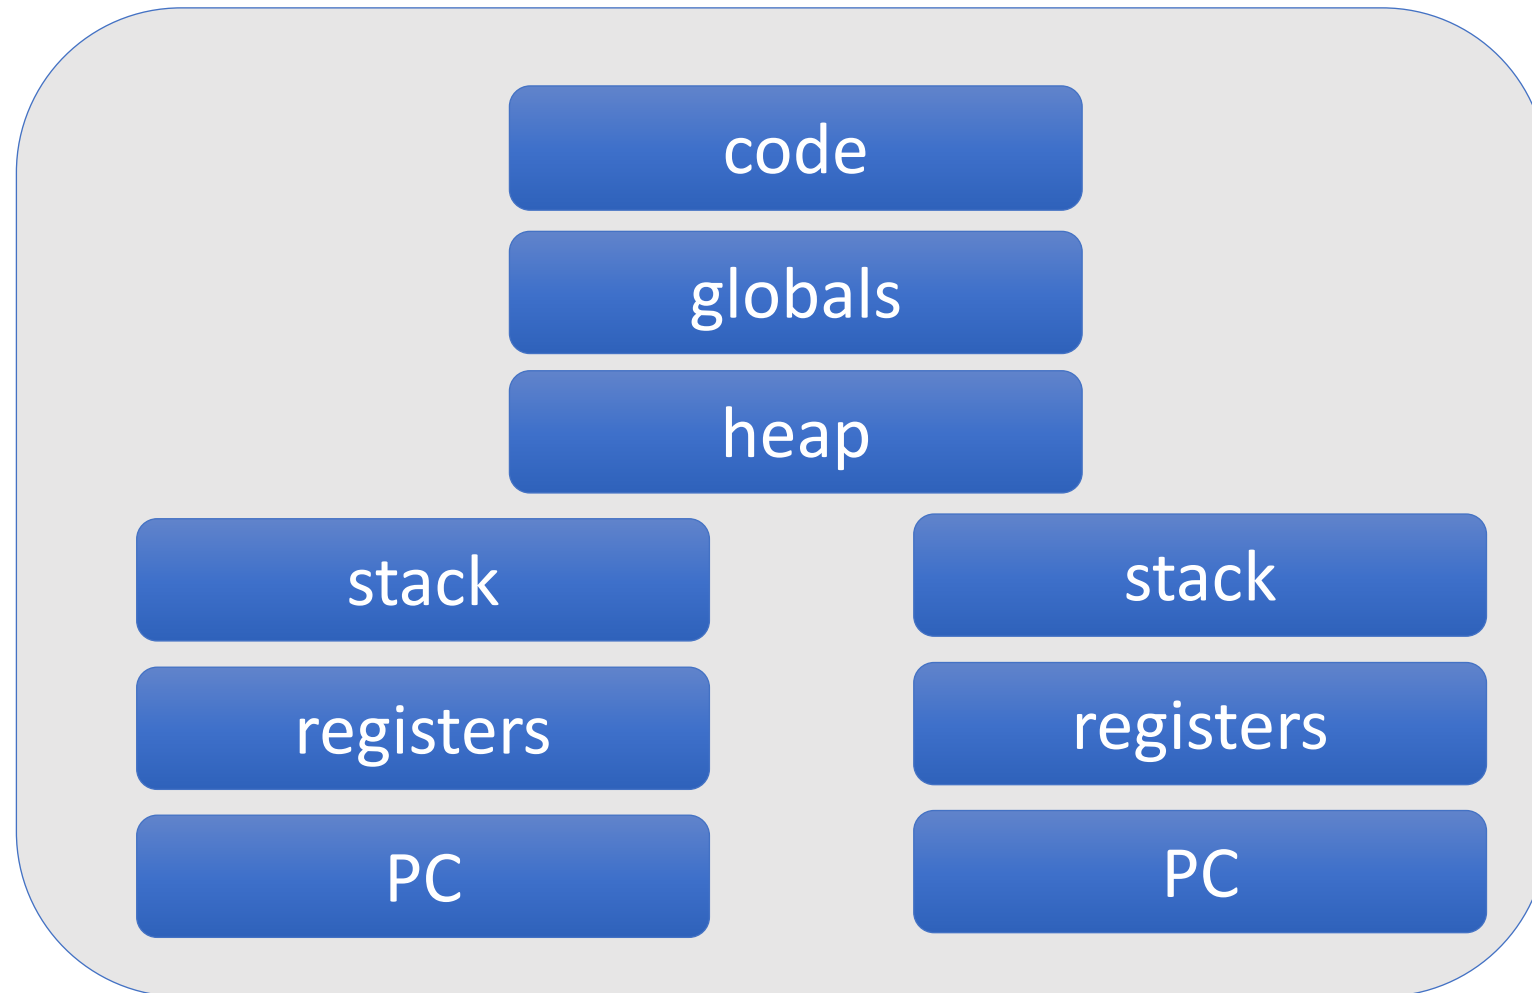
Will see synchronization principles **today**

Will see practical examples later

Two Processes




Two Threads in a Process



In General

- Processes provide separation
 - In particular, memory separation (no shared data)
 - Suitable for coarse-grain interaction
- Threads do not provide separation
 - In particular, threads share memory (shared data)
 - Suitable for tighter integration

Where Things Reside



Shared Data

- Advantage:
 - Many threads can read/write it
- Disadvantage:
 - Many threads can read/write it

Shared Data

- Advantage:
 - Many threads can read/write it
- Disadvantage:
 - Many threads can read/write it
 - Can lead to *data races*

Data Race

- Unexpected/unwanted access to shared data
- Result of *interleaving* of thread executions
- **Program must be correct for all interleavings**
 - (another COMP 525 topic)

A Common Mistake/Misunderstanding: A Single Line of Code is not Atomic

- $a = a + 1$
- Is in reality
 - Load a from memory into register
 - Increment register
 - Store register value in memory
- Instruction sequence may be interleaved
- (Some machines have atomic increments)

Thread Schedule #1

```
balance = balance + 1; // balance in shared memory at 0x9cd4
```


Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 100

%eax: ?

%rip = 0x195

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 100

%eax: 101

%rip = 0x19d

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 →

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195

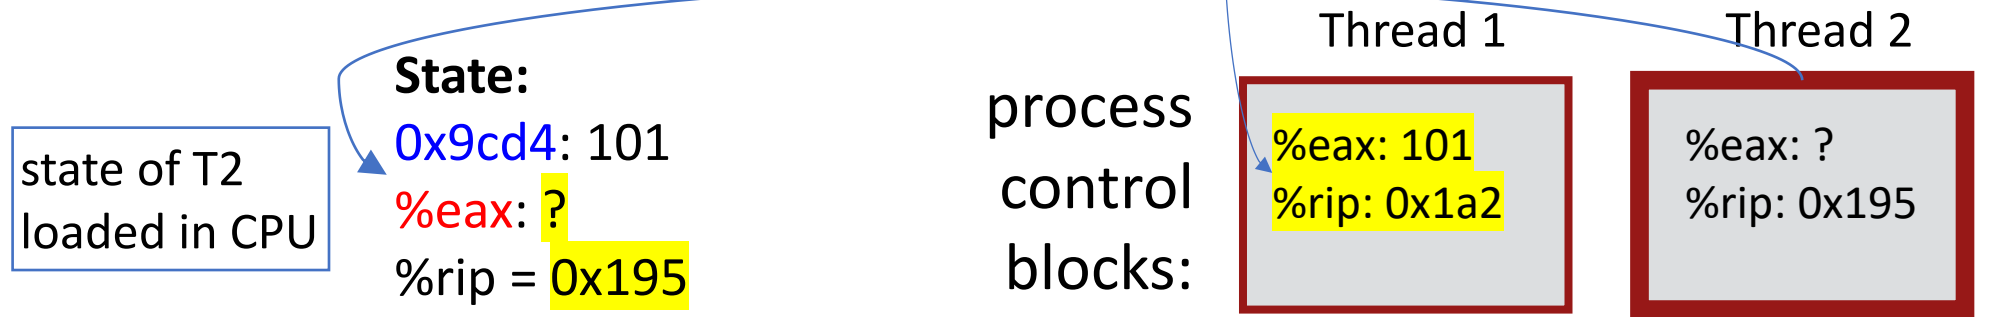
- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4



Thread context switch!

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

note that code region is common to T1 & T2

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 101

%eax: ?

%rip = 0x195


process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x19a


process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 101

%eax: 102

%rip = 0x19d


process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #1

balance = balance + 1; // balance in shared memory at 0x9cd4

Desired
result 😊

State:

0x9cd4: 102

%eax: 102

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x1a2

Thread 2

%eax: ?
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T2 →

Another schedule

Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 100

%eax: ?

%rip = 0x195

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a


process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 100

%eax: 101

%rip = 0x19d

process
control
blocks:

Thread 1

%eax: ?
%rip: 0x195

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread context switch!

Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 100

%eax: ?

%rip = 0x195


process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

T2 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 100

%eax: 101

%rip = 0x19d


process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T2



Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: ?
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T2



Thread context switch!

Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

State:

0x9cd4: 101

%eax: 101

%rip = 0x19d

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: 101
%rip: 0x1a2



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

Thread Schedule #2

balance = balance + 1; // balance in shared memory at 0x9cd4

WRONG Result! ☹️
**Final value of
balance is 101**

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process
control
blocks:

Thread 1

%eax: 101
%rip: 0x19d

Thread 2

%eax: 101
%rip: 0x1a2

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 →

Non-Determinism

Concurrency leads to non-deterministic results

- Different results even with same inputs
- Race conditions

Whether bug manifests or not depends on CPU schedule!

How to program?

- **Assume scheduler is malicious**
- Assume scheduler will pick bad ordering at some point...

Basic Approach to Multithreading

1. Divide “work” among multiple threads &
2. Share data
 - Which data is shared?
 - **Global variables and heap**
 - Not local variables, not read-only variables
 - Where is shared data accessed?
 - Put shared data access in **critical section**

Critical Section

- Want 3 instructions to execute as an uninterruptable group
- We say we want them to be “**atomic**”

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

critical section

Need **mutual exclusion** for critical sections

- If thread A is in critical section C, thread B can't enter C
- Ok if other processes/threads do unrelated work

Aside: Non-Atomic Single Instructions

- We just implemented `x = x + 1` with:

```
mov 0x123, %eax  
add $0x1, %eax  
mov %eax, 0x123
```

- But x86 has this instruction:

```
add $0x1, 0x123
```

Aside: Non-Atomic Single Instructions

- We just implemented ``x = x + 1`` with:

```
mov 0x123, %eax
add $0x1, %eax
mov %eax, 0x123
```

- But x86 has this instruction:

```
add $0x1, 0x123
```

- No difference! **This single instruction is not atomic.**
 - (The similar ``lock add`` instruction is atomic.)

Mutual Exclusion

- Prevents simultaneous access to a shared resource.
 - In this case, shared resource = shared memory region
- How can we achieve mutual exclusion?
 - Today we will first see library support (pthreads)
 - Then, we will see implementation of synchronization primitives.

Why this (mostly) works

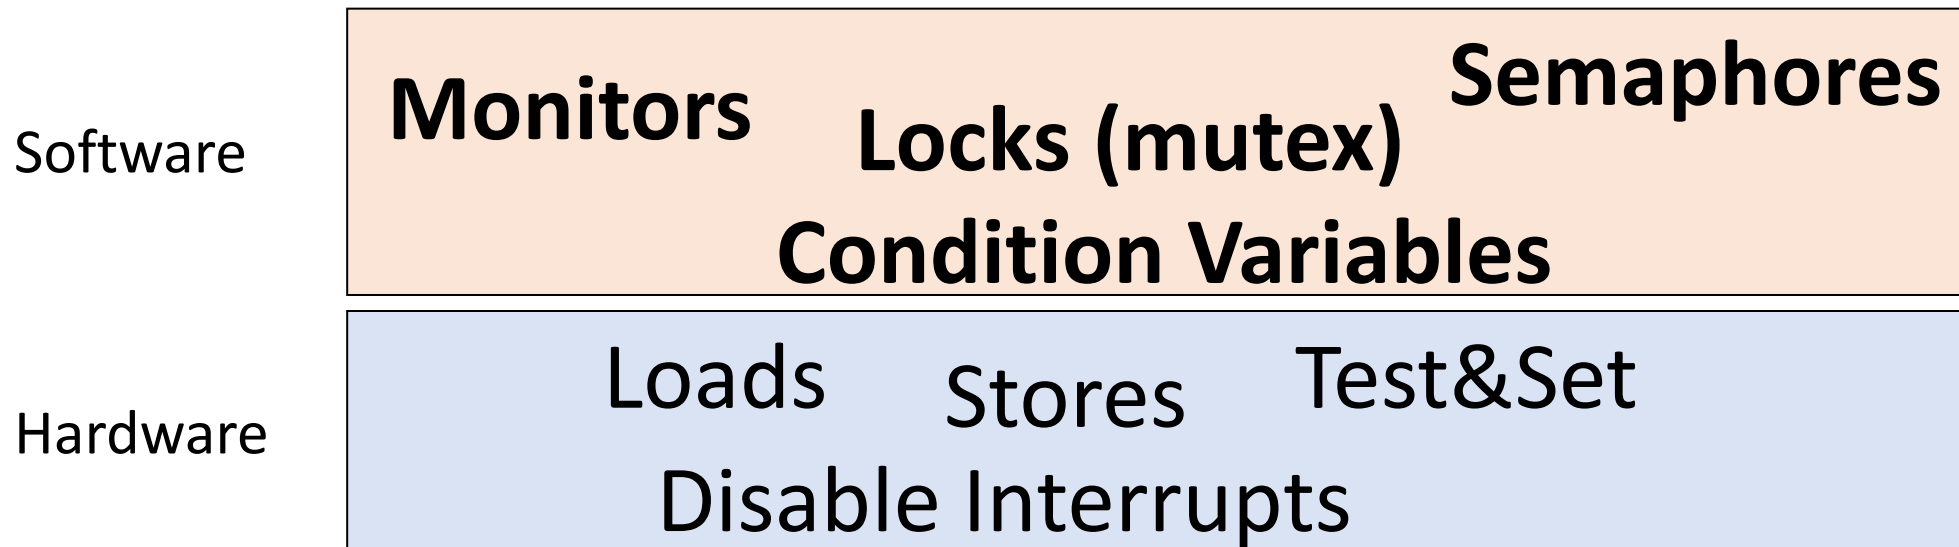
- Critical section:
 - No other thread can change data
- So you are (mostly) ok

Synchronization

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



POSIX Thread Libraries (pthreads)

- Thread API for C/C++
- User-level library:
 - **#include <pthread.h>**
 - Compile and link with *-pthread*.
- Support for thread creation, termination, synchronization.
- See more details here: <https://man7.org/linux/man-pages/man3/>

Pthreads: Thread Creation and Destruction

- `pthread_create()`
- `pthread_exit()`
- `pthread_join()`

pthread_create()

```
int pthread_create(  
    pthread_t * thread, pthread_attr_t * attr,  
    void *(*start_routine)(void *), void * arg  
);
```

- Create thread, in *thread*.
- Run *start_routine* with arguments *arg*.
- *attr* points to a *pthread_attr_t* structure. If *attr* is NULL, then the thread is created with default attributes (ok in most cases).
- On success, return 0; on error, return an error number.

pthread_exit()

```
void pthread_exit(void *retval);
```


- Terminate calling thread.
- Returns a value via *retval*.

pthread_join()

```
int pthread_join(pthread_t thread, void **retval);
```

- Join with a terminated thread.
- Waits for the thread specified by *thread* to exit.
- If *retval* is not NULL, then **pthread_join()** copies the return value of the target thread into the location pointed to by *retval*.
- On success, return 0; on error, return an error number.

Fork-Join Pattern for threads



Main thread creates (forks) collection of sub-threads passing them args to work on...
... and then joins with them, collecting the results

Note: In this example, start_routine is the same for all threads; only args are different.

Fork-join example

```
void *mythread(void *arg) {  
    printf("%s\n", (char *) arg);  
    return NULL;  
}  
  
int main(int argc, char *argv[]) {  
    pthread_t p1, p2;  
    printf("main: begin\n");  
    pthread_create(&p1, NULL, mythread, "A");  
    pthread_create(&p2, NULL, mythread, "B");  
    // join waits for the threads to finish  
    pthread_join(p1, NULL);  
    pthread_join(p2, NULL);  
    printf("main: end\n");  
}
```

Counting example – What is the final answer?

```
int count;
void *mythread(void *arg) {
    int j;
    for (j = 0; j < 1000000; j++){
        count +=1;
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p1, p2;
    count = 0;
    pthread_create(&p1, NULL, mythread, NULL);
    pthread_create(&p2, NULL, mythread, NULL);
    pthread_join(p1, NULL);
    pthread_join(p2, NULL);
    printf("%d \n", count);
}
```

Pthreads: Locks

- `Pthread_mutex_lock(mutex)`
- `Pthread_mutex_unlock(mutex)`

Pthread_mutex_lock(mutex)

- If lock is held by another thread, block
- If lock is not held by another thread
 - Acquire lock
 - Proceed

Pthread_mutex_unlock(mutex)

- Release lock

Counting example revisited – What is the final answer?

```
pthread_mutex_t count_mutex;
int count;

void *mythread(void *arg) {
    int j;
    for (j = 0; j < 1000000; j++){
        pthread_mutex_lock(&count_mutex);
        count +=1;
        pthread_mutex_unlock(&count_mutex);
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p1, p2;
    pthread_mutex_init(&count_mutex, NULL);

    count = 0;
    pthread_create(&p1, NULL, mythread, NULL);
    pthread_create(&p2, NULL, mythread, NULL);
    pthread_join(p1, NULL);
    pthread_join(p2, NULL);
    printf("%d \n", count);
}
```

Deadlocks

- Threads are stuck waiting for blocked resources and no amount of retry (backoff) will help.
- Classic example:

Thread A

```
1 lock(object1)
2 lock(object2)
3 //do stuff
4 unlock(object2)
5 unlock(object1)
...
```

Thread B

```
1 lock(object2)
2 lock(object1)
3 //do stuff
4 unlock(object1)
5 unlock(object2)
...
```



Deadlocks


- Threads are stuck waiting for blocked resources and no amount of retry (backoff) will help.
- Classic example:

Thread A

```
1 lock(object1)
2 lock(object2)
3 //do stuff
4 unlock(object2)
5 unlock(object1)
...
```

Thread B

```
1 lock(object2)
2 lock(object1)
3 //do stuff
4 unlock(object1)
5 unlock(object2)
...
```



Deadlocks


- Threads are stuck waiting for blocked resources and no amount of retry (backoff) will help.
- Classic example:

Thread A

```
1 lock(object1)
2 lock(object2)
3 //do stuff
4 unlock(object2)
5 unlock(object1)
...
```

Thread B

```
1 lock(object2)
2 lock(object1)
3 //do stuff
4 unlock(object1)
5 unlock(object2)
...
```



Deadlock example

```
pthread_mutex_t lock1;
pthread_mutex_t lock2;

void *a_func(void *arg) {
    long j;
    for (j = 0; j < 1000000000; j++) {
        pthread_mutex_lock(&lock1);
        pthread_mutex_lock(&lock2);
        printf("A");
        pthread_mutex_unlock(&lock2);
        pthread_mutex_unlock(&lock1);
    }
    return NULL;
}
```

```
void *b_func(void *arg) {
    long j;
    for (j = 0; j < 1000000000; j++) {
        pthread_mutex_lock(&lock2);
        pthread_mutex_lock(&lock1);
        printf("B");
        pthread_mutex_unlock(&lock1);
        pthread_mutex_unlock(&lock2);
    }
    return NULL;
}
```

```
int main(int argc, char *argv[]) {
    pthread_t a, b;
    pthread_mutex_init(&lock1, NULL);
    pthread_mutex_init(&lock2, NULL);
    pthread_create(&a, NULL, a_func, NULL);
    pthread_create(&b, NULL, b_func, NULL);
    pthread_join(a, NULL);
    pthread_join(b, NULL);
    printf("End!\n");
}
```

Week 3

Synchronization Primitives

Max Kopinsky
23 January, 2025

Assignment 1 Release – Clone

- Clone your fork
 - on mimi
 - (+ your machine, if you want)
 - Use your SOCS user+passwd

This command (with your own repo link)

```
mkopin@teach-node-04:~  
$ git clone https://gitlab.cs.mcgill.ca/mkopin/operating-systems-w25  
Cloning into 'operating-systems-w25'...  
Username for 'https://gitlab.cs.mcgill.ca': mkopin  
Password for 'https://mkopin@gitlab.cs.mcgill.ca':  
warning: redirecting to https://gitlab.cs.mcgill.ca/mkopin/operating-systems-w25.git/  
remote: Enumerating objects: 6, done.  
remote: Counting objects: 100% (6/6), done.  
remote: Compressing objects: 100% (4/4), done.  
remote: Total 6 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)  
Receiving objects: 100% (6/6), done.
```

either of these

Clone with SSH
git@gitlab.cs.mcgill.ca:mkopin/o

Clone with HTTPS
mkopin/operating-systems-w25.git

Open in your IDE
Visual Studio Code (SSH)

Name	Last commit
README.md	initial commit


You can ignore this warning if you get it

Assignment 1 Release – Visibility Settings

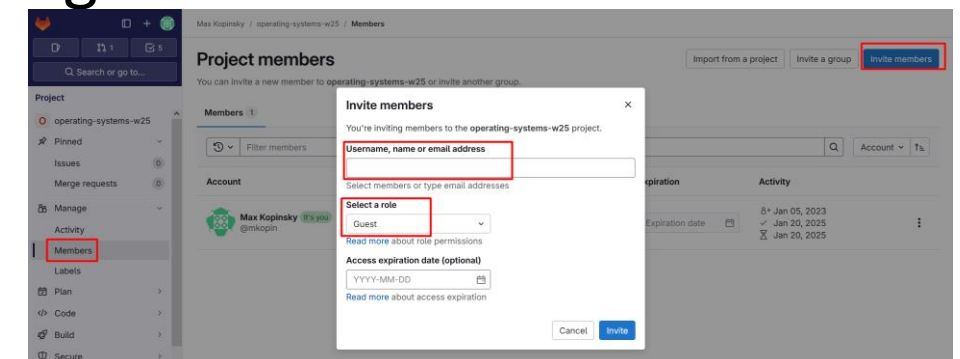
- From the GitLab page for your repository
- Left panel > settings > General > Visibility
 - Change “Project Visibility” to **private**
 - If you do not do this, other students can copy your work. You would **both** be in trouble.

✓ Visibility, project features, permissions

Choose visibility level, enable/disable project features and their permissions, disable email notifications, and show default emoji reactions.



- Add your partner, myself, and the autograding TA to members:
 - Partner role: developer
 - Also add: @mkopin (me!) with role Reporter
 - And add @mohamad.danesh as Reporter too



Assignment 1 Release – Git Merge

- Finally whenever we release an assignment, you must update your fork with the changes from our “upstream” fork to get the materials (instructions, test cases, code, ...) for that assignment.
- These commands will also be posted on Discord.

```
$ git remote add staff https://gitlab.cs.mcgill.ca/mkopin/operating-systems-w25  
$ git fetch staff  
$ git checkout main  
$ git merge staff/main
```

This one must be to my repository

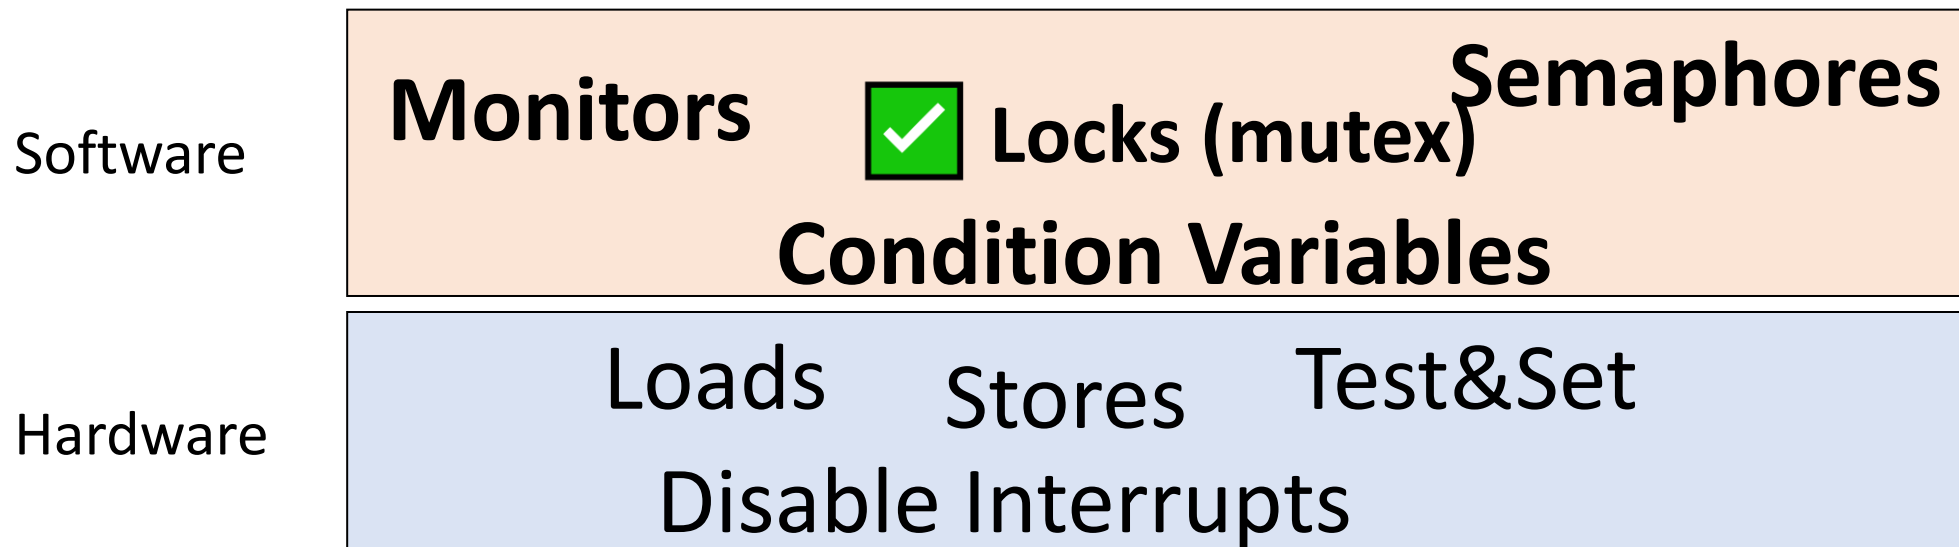


Synchronization

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



Condition Variables

- Used when thread A needs to wait for an event done by thread B
- A waits until a certain condition is true
 - First test condition,
 - If condition not true, call `pthread_cond_wait()`
 - A blocks until another thread “signals” the condition variable
- At some point B makes the condition true
 - Then B calls `pthread_cond_signal()`, which unblocks A.

Advantage of Condition Variable of Locks

- A waits until a certain condition is true
 - A goes to sleep here.
 - Does not keep CPU busy.

Condition variables use (incorrect)

Thread A

```
x = f (a , b) ;  
if (x < 0 || x > 9)  
    pthread_cond_wait (&cv);
```

Thread B

```
//change a and b;  
x = f (a , b) ;  
if ( x >= 0 && x <= 9)  
    pthread_cond_signal (&cv);
```

Find the data race.

Condition variables use (incorrect)

Thread A

```
x = f ( a , b );  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv);
```

Interrupt

Thread B

```
//change a and b;  
x = f ( a , b );  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv);
```

Condition variables use (incorrect)

Thread A

```
x = f ( a , b );  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv);
```

Interrupt

Thread B

```
//change a and b;  
x = f ( a , b );  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv);
```

Condition variables use (incorrect)

Thread A

```
x = f ( a , b );  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv);
```

Interrupt

Thread B

```
//change a and b;  
x = f ( a , b );  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal(&cv);
```

:(Broadcast missed by A

Condition variables use (incorrect)

Thread A

```
x = f ( a , b );  
if ( x < 0 || x > 9 )  
    pthread_cond_wait (&cv);
```

A waits forever...

Thread B

```
//change a and b;  
x = f ( a , b );  
if ( x >= 0 && x <= 9 )  
    pthread_cond_signal (&cv);
```

Condition variables use (still incorrect)

Thread A

```
pthread_mutex_lock(&mutex);  
x = f(a, b);  
if (x < 0 || x > 9)  
    pthread_cond_wait(&cv, &mutex);  
pthread_mutex_unlock(&mutex);
```

Thread B

```
pthread_mutex_lock(&mutex);  
//change a and b;  
x = f(a, b);  
if (x >= 0 && x <= 9)  
    pthread_cond_signal(&cv);  
pthread_mutex_unlock(&mutex);
```

Remember: Condition variable is a shared resource between A and B
→ Every time you use a condition variable you must also use a mutex to prevent the race condition.

One more issue...

Sometimes, the wait function might return even though the condition variable has not actually been signaled. **This is part of the API** as an unfortunate necessity.

Thread A

```
pthread_mutex_lock(&mutex);  
x = f(a, b);  
if (x < 0 || x > 9)  
    pthread_cond_wait(&cv, &mutex);  
pthread_mutex_unlock(&mutex);
```

Thread B

```
pthread_mutex_lock(&mutex);  
//change a and b;  
x = f(a, b);  
if (x >= 0 && x <= 9)  
    pthread_cond_signal(&cv);  
pthread_mutex_unlock(&mutex);
```

Example:

If process P running A and B receives an OS signal

- Any thread in P can be chosen to process the signal.
- → A might be chosen to process the signal handling function
- → wait returns with an error code → A runs even if condition is not true...

How can we fix this?

Condition variables use (correct)

- Retest the condition after `pthread_cond_wait()` returns.
 - This is most easily done using a loop.

Thread A

```
pthread_mutex_lock(&mutex);  
while (1) {  
    x = f(a, b);  
    if (x < 0 || x > 9) {  
        pthread_cond_wait(&cv, &mutex);  
    } else {  
        break;  
    }  
}  
pthread_mutex_unlock(&mutex);
```

Thread B

```
pthread_mutex_lock(&mutex);  
//change a and b;  
x = f(a, b);  
if (x >= 0 && x <= 9)  
    pthread_cond_signal(&cv);  
pthread_mutex_unlock(&mutex);
```

Conditional Variables Interface

- `pthread_cond_init(pthread_cond_t *cv, pthread_condattr_t *cattr)`
 - Initialize the conditional variable, `cattr` can be NULL
- `pthread_cond_wait(pthread_cond_t *cv, pthread_mutex_t *mutex)`
 - Block thread until condition is true, and atomically unlock mutex.
 - When the thread is unblocked, reacquire the mutex as if by `pthread_mutex_lock`
 - Might cause thread to block again!
- `pthread_cond_signal(pthread_cond_t *cv)`
 - Unblock one thread at random that is blocked by the condition variable
- `pthread_cond_broadcast(pthread_cond_t *cv)`
 - Unblock all threads that are blocked on the condition variable pointed to by `cv`.

Condition Variable Example

```
pthread_cond_t is_zero;
pthread_mutex_t mutex;
int shared_data = 100;

void *thread_func(void *arg){
    while(shared_data > 0) {
        pthread_mutex_lock(&mutex);
        shared_data--;
        printf("%d ", shared_data);
        pthread_mutex_unlock(&mutex);
    }

    printf("Signaling main\n");
    pthread_cond_signal(&is_zero);
    return NULL;
}
```

```
int main (void){

    pthread_t tid;
    void * exit_status;
    int i;

    pthread_cond_init(&is_zero, NULL);
    pthread_mutex_init(&mutex, NULL);

    pthread_create(&tid, NULL, thread_func, NULL);

    pthread_mutex_lock(&mutex);
    printf("Start waiting in main\n");
    while(shared_data!=0)
        pthread_cond_wait(&is_zero, &mutex);
    pthread_mutex_unlock(&mutex);

    printf("Done waiting in main!\n");


    pthread_join(tid, &exit_status);
    return 0;
}
```

Synchronization


Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



Semaphores



What are Semaphores?


- A shared, non-negative counter.
- Two primary operations:
 - Wait → attempts to decrement the counter; blocks when counter is 0.
 - Post (or Signal) → attempts to increment the counter.
- `#include<semaphore.h>`

Changes are atomic



unsigned int

Semaphore




Semaphore

Changes are atomic
2 operations

wait()

post()



Semaphore


Changes are atomic
2 operations

wait()

try to decrement semaphore value
block if value = 0

post()

increment semaphore value



Semaphore

wait() {

while(1){

if (v>0){

v--;

return;


}

}

}

Atomic execution

Note: wait() is not actually implemented like this. This is how the behavior looks like to the programmer



Semaphore

wait() {

while(1){

if (v>0){

v--;

return;


}

}

}

Atomic execution

Note: wait() is not actually implemented like this. This is how the behavior looks like to the programmer



Semaphore

wait() {

while(1){

if (v>0){

v--;

return;

}


}

}

Atomic execution

Wait until semaphore value becomes positive again

Note: wait() is not actually implemented like this. This is how the behavior looks like to the programmer



Semaphore

wait() {

while(1){

if (v>0){

v--;

return;

}


}

}

Atomic execution

Once the value is positive, the thread that
Was waiting is able to decrement the value.

Note: wait() is not actually implemented like this. This is how the behavior looks like to the programmer




Semaphore

post() {

```
v++;  
return;  
}
```

Atomic execution

Note: post() is not actually implemented like this. This is how the behavior looks like to the programmer



Semaphore

post() {

```
v++;  
return;  
}
```

Atomic execution

Note: post() is not actually implemented like this. This is how the behavior looks like to the programmer



unsigned int

Semaphore

One last thing:

- Semaphore value can be initialized upon creation to a positive value
- or 0

Semaphore Uses

Mutual exclusion

- A semaphore with its counter initialized to 1 can be used as a lock.

Bound the concurrency

- Only allow X threads out of N to proceed.

Producer-consumer problem

- More complex use of semaphores. Will see in 2 weeks.

Semaphores interface

- `int sem_init(sem_t *sem, int pshared, unsigned value);`
- `int sem_post(sem_t *sem);`
- `int sem_wait(sem_t *sem);`

For more details: <https://man7.org/linux/man-pages/man0/semaphore.h.0p.html>

```
int sem_init(sem_t *sem, int pshared, unsigned value);
```

- Initializes the semaphore * *sem*. The initial value of the semaphore is *value*.
- If *pshared* is 0, the semaphore is shared among all threads of a process.
- If *pshared* is not zero, the semaphore is shared *between processes* but
 - Must be stored somewhere that multiple processes can see, e.g. file-mapped memory.
 - Don't worry about this use case. We won't discuss process-shared memory in this course.
- Return 0 on success, -1 on failure.

```
int sem_wait(sem_t *sem);
```

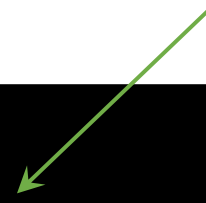
- If the *sem* has a value > 0 , decrement the value by 1.
- If *sem* has value 0, the caller will be blocked until *sem* has a value larger than 0.
- Return 0 on success, -1 on failure.

```
int sem_post(sem_t *sem);
```

- Increment the value of *sem* by 1.
- If threads are blocked waiting for the semaphore, one of them (at random) will return successfully from its call to *sem_wait()*; the semaphore value is immediately decremented.
 - (this doesn't happen until one of those threads is next scheduled – but focus on concepts!)
- Return 0 on success, -1 on failure.

Semaphores example

Note: If 2 is changed to 1, we have a lock behavior



```
#include <pthread.h>
#include <semaphore.h>

pthread_t threads[5];
int tid[5];
sem_t sem;

void * thread_func(void *arg){
    int tid_ = tid[* (int *) arg];
    printf("Thread %d created\n", tid_);
    int j;

    sem_wait(&sem);
    for (j=0; j<3; j++){
        printf("T%d run %d\n", tid_, j);
        sleep(2);
    }
    sem_post(&sem);
}
```

```
int main(){

    sem_init(&sem, 0, 2);
    //sem initialized for all threads in the process; allow only 2 threads in
    the critical section at a time;

    int i;
    for (i=0; i<5; i++){
        tid[i]=i;
        pthread_create(&threads[i], NULL, thread_func, &tid[i]);
    }

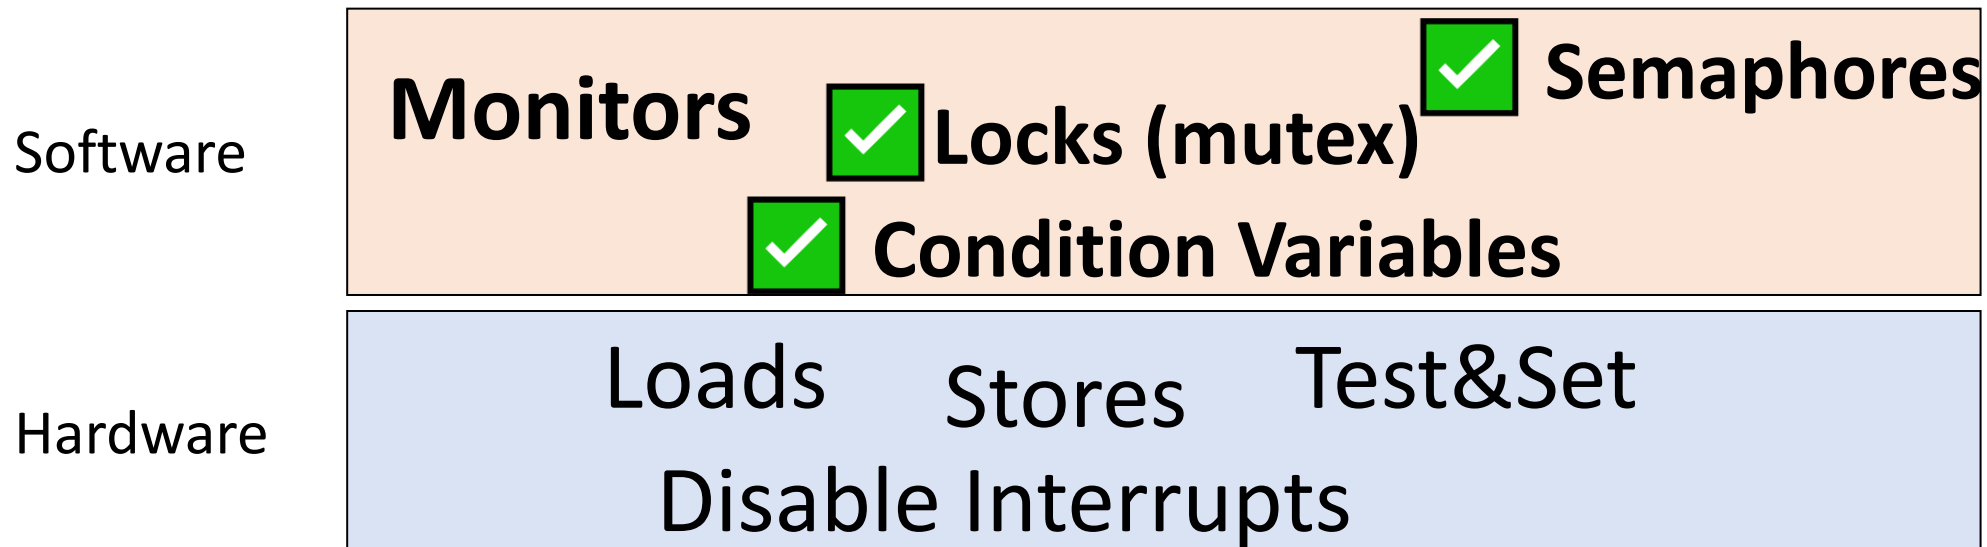
    for (i=0; i<5; i++){
        pthread_join(threads[i], NULL);
    }
    sem_destroy(&sem);
    return 0;
}
```

Synchronization

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



Monitors

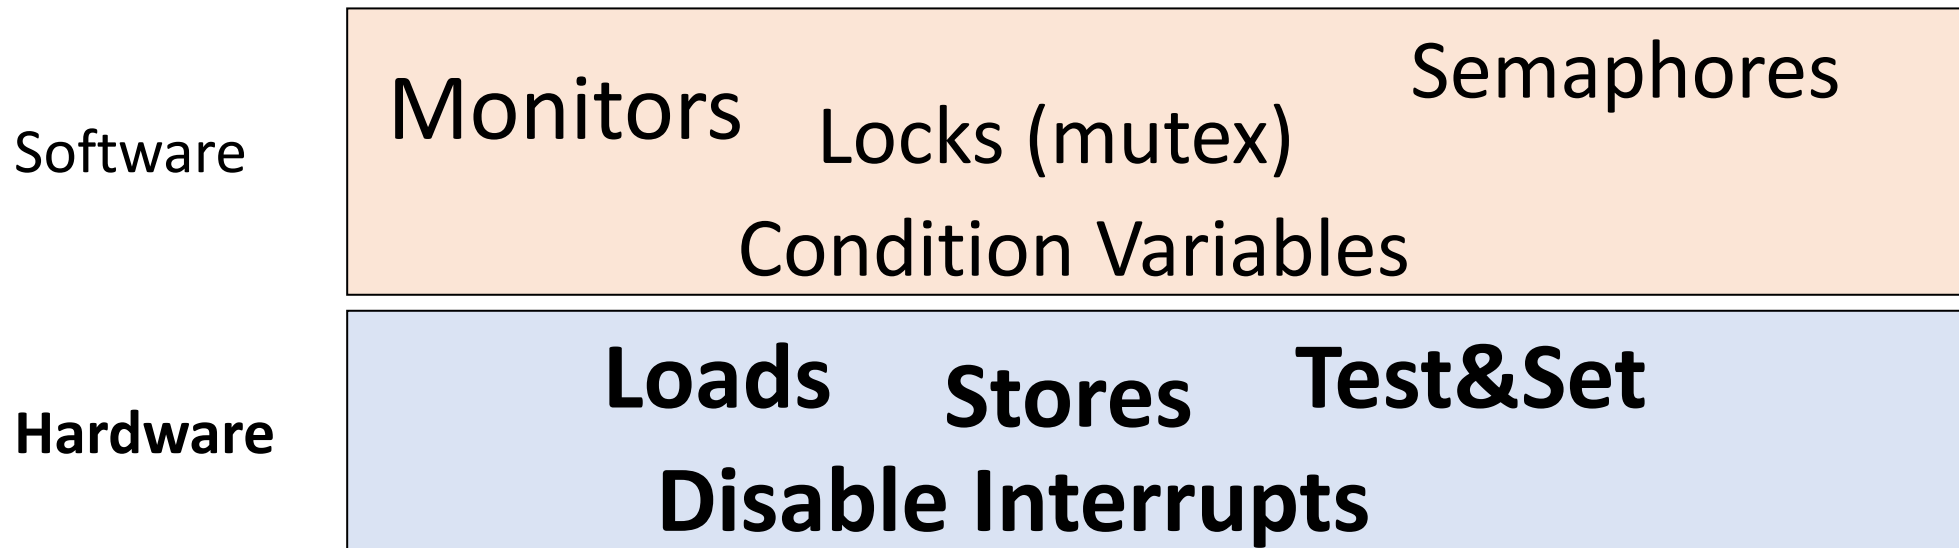
- Collection of variables and functions
- Threads can only access monitor functions
 - Variables are private to the monitor
- Only one process at a time can execute code inside the monitor.

Monitors

- Pthreads does not offer a monitor primitive 😞
- ... but possible to implement monitor semantics using mutex and condition variables 😊
- We won't get into this in this course. Investigate if you're curious! 😊

Synchronization

How are locks, semaphores, cond var, etc **implemented?**



Reminder of Assumptions:

- Multitasking OS
 - “Tasks” can be threads or processes
- Single-processor system
 - Only one task runs at any time
- **Not** assuming OS helps with synchronization
 - Of course, all the modern OSes do
 - But if you need to know how to implement... its because you don't have one.

Lock implementation motivating example

“Too much milk”

- Alice and Bob are roommates
- They want to coordinate grocery shopping
 - Need to be careful to not buy too much of perishable items, like milk.



Schedule that leads to too much milk

Time	Alice	Bob
3:00	Look in Fridge. Out of milk.	
3:05	Leave for store.	
3:10	Arrive at store.	Look in Fridge. Out of milk.
3:15	Buy milk.	Leave for store.
3:20	Arrive home, put milk away.	Arrive at store.
3:25		Buy milk.
3:30		Arrive home, put milk away.

Problem specifications

Safety

- Never more than one person buys

Liveness

- Someone buys if needed

Lock implementation: first attempt

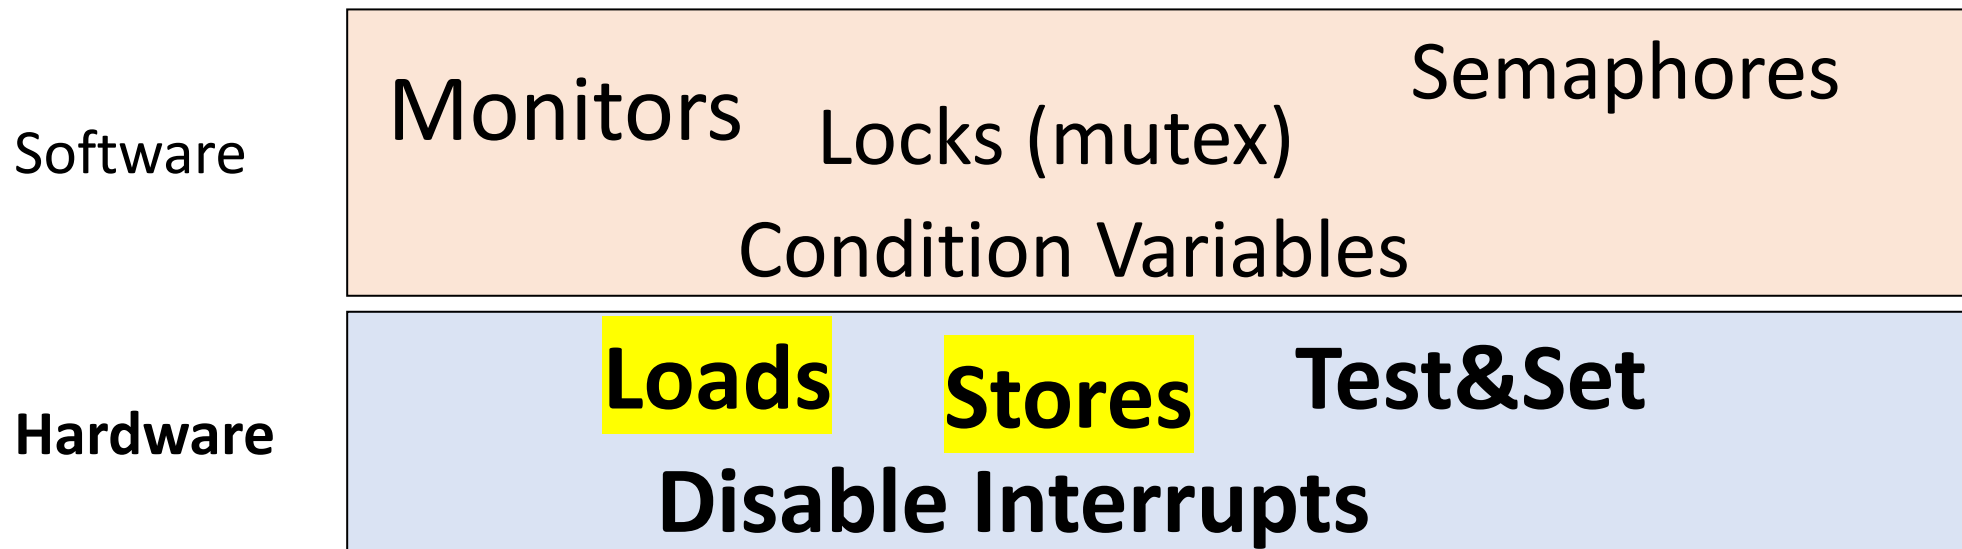
Restrict ourselves to use only atomic load and store operations as building blocks.

Idea: Use **a note** to avoid buying too much milk:

- Leave a note on fridge before buying (kind of “lock”)
- Remove note after buying (kind of “unlock”)
- Don’t buy if note (wait)

Synchronization

How are locks, semaphores, cond var, etc **implemented?**



Lock implementation: first attempt (**incorrect**)

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



Lock implementation: first attempt (incorrect)

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note; }  
}
```

load (atomic)

store (atomic)



Lock implementation: first attempt (incorrect)

Thread A

```
if (noMilk) {
```

Thread B

```
if (noMilk) {  
    if (noNote) {
```

Aside: when I want to anthropomorphize threads/processes like this, thread A will be “Alice” and thread B will be “Bob.”

Lock implementation: first attempt (incorrect)

Thread A

```
if (noMilk) {
```

```
    if (noNote) {
```

```
        leave Note;
```

```
        buy milk;
```

```
        remove note;
```

```
    }
```

```
}
```

Thread B

```
if (noMilk) {
```

```
    if (noNote) {
```

Lock implementation: first attempt (incorrect)

Thread A

```
if (noMilk) {
```

```
    if (noNote) {
```

```
        leave Note;
```

```
        buy milk;
```

```
        remove note;
```

```
    }
```

```
}
```

Buy milk twice 😞

Remember: scheduler can create any interleaving

We must assume a malicious scheduler

Thread B

```
if (noMilk) {
```

```
    if (noNote) {
```

```
        leave Note;
```

```
        buy milk;
```

```
        remove note; }
```

```
}
```


First attempt result

- Still too much milk but only occasionally!
 - This is worse than a consistent error, because it is harder to catch.
- Thread can get context switched **after checking milk** (and note) **but before buying** milk!

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

```
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note B;
```

Problem solved?

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

Proc switch leave note A;
if (noNote B) {
 if (noMilk) {
 buy milk;
 }
}
remove note A;

Thread B

leave note B;
if (noNote A) {
 if (noMilk) {
 buy milk;
 }
}
remove note B;

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

leave note A;

if (noNote B) {

if (noMilk) {

buy milk;

}

}

remove note A;

Proc switch

Thread B

leave note B;

if (noNote A) {

if (noMilk) {

buy milk;

}

}

remove note B;

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

leave note A;

if (noNote B) {

if (noMilk) {

buy milk;

}

}

remove note A;

Thread B

leave note B;

if (noNote A) {

if (noMilk) {

buy milk;

}

}

remove note B;

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

leave note A;

if (noNote B) {

if (noMilk) {

buy milk;

}

}

Proc switch

remove note A;

Thread B

leave note B;

if (noNote A) {

if (noMilk) {

buy milk;

}

}

remove note B;

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

leave note A;

if (noNote B) {

if (noMilk) {

buy milk;

}

}

remove note A;

Thread B

leave note B;

if (noNote A) {

if (noMilk) {

buy milk;

}

}

remove note B;

Proc switch

Not quite...

Lock implementation: second attempt (incorrect)

- How about labeled notes?

Thread A

leave note A;

if (noNote B) {

if (noMilk) {

buy milk;

}

}

remove note A;

Thread B

leave note B;

if (noNote A) {

if (noMilk) {

buy milk;

}

}

remove note B;

Not quite...

Nobody buys milk 😞

Lock implementation: third attempt (correct)

- How about labeled notes? (similar to Peterson, 1981)

Thread A

leave note A;

while (note B)

do nothing;

if (noMilk)

buy milk;

remove note A;

Thread B

leave note B;

if (noNote A) {

if (noMilk)

buy milk;

}

remove note B;

This works!

B has priority to buy

Correctness argument. Case 1

Thread A

```
leave note A;  
while (note B)  
  do nothing;  
if (noMilk)  
  buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
  if (noMilk)  
    buy milk;  
}  
remove note B;
```

Happened before



Correctness argument. Case 1

Thread A

```
leave note A;  
while (note B)  
do nothing;
```

```
if (noMilk)  
buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before



Correctness argument. Case 1

Thread A

```
leave note A;  
while (note B)  
do nothing;
```



Wait for note B
to be removed

```
if (noMilk)  
buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before



Correctness argument. Case 1

Thread A

```
leave note A;  
while (note B)  
do nothing;
```

Wait for note B
to be removed

```
if (noMilk)  
buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before



Correctness argument. Case 2

Thread A

```
leave note A;  
while (note B)  
  do nothing;  
if (noMilk)  
  buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
  if (noMilk)  
    buy milk;  
}  
remove note B;
```

Happened before


Correctness argument. Case 2

Thread A

```
leave note A;  
while (note B)  
do nothing;  
if (noMilk)  
    buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
    if (noMilk)  
        buy milk;  
}  
remove note B;
```

Happened before


Correctness argument. Case 2

Thread A

```
leave note A;  
while (note B)  
do nothing;
```

```
if (noMilk)  
buy milk;  
remove note A;
```

Thread B

```
leave note B;  
if (noNote A) {  
  if (noMilk)  
    buy milk;  
}  
remove note B;
```

Happened before

Wait for note B
to be removed

Prove Correctness?

- By hand, surprisingly hard and often requires insight
- Mechanical tool exists: “Linear Temporal Logic”
- I can show an LTL specification + proof of solution 3 in OH if interested
- Even with LTL, still surprisingly tricky!
- If this kind of automated proof is interesting, try take 525
 - (Hopefully, it’s taught again soon 😞)

Attempt 3 discussion

Solution 3 works, but it's really unsatisfactory.

- Complex, even for this simple example.
 - Hard to convince yourself that this really works.
- What we have guarantees that *A or B* can enter critical section
 - But we usually want to know that they *both* can, eventually
 - (In this case, we just know that they are trying to do the same thing)
- A's code is different from B's –what if lots of threads?
 - Code would have to be slightly different for each thread.
- While A is waiting, it is consuming CPU time.
 - This is called “busy-waiting”.

There must be a better way!

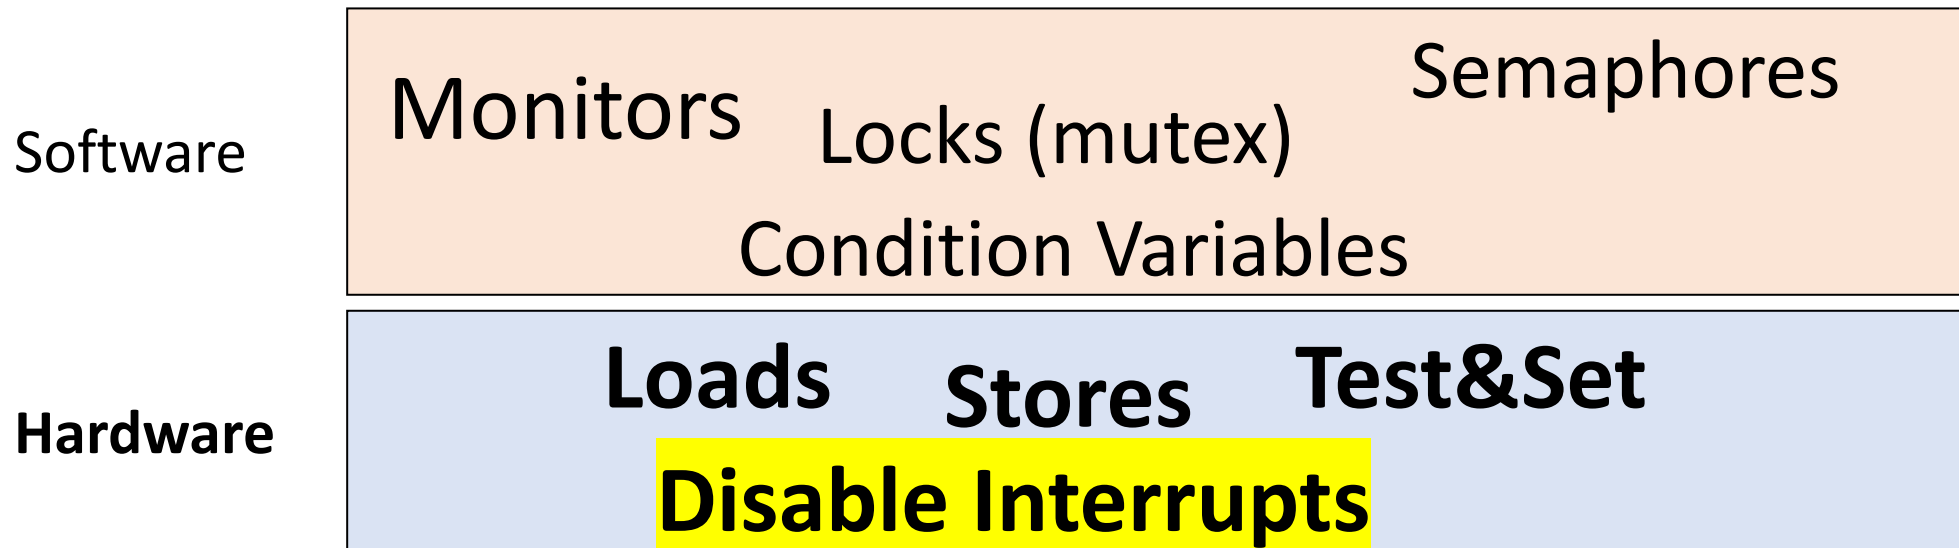
- Version of Solution 3 where both threads have same code exists
 - See Peterson's Algorithm (1981)
- Have higher-level hardware primitives than atomic load & store
 - Compare-and-swap, test&set, LL/SC
- Build higher-level abstractions on this hardware support
- **Or: implement abstractions at OS level**

Fun Fact: It Gets Worse

- Modern multicore processors execute instructions *out of order*
- Normally not possible to notice
 - Hardware is carefully designed so that single thread can't tell
- But multiple threads can observe re-ordered memory accesses
- **Re-ordering note bookkeeping breaks Solution 3**
- With only atomic loads/stores (and nothing “stronger”), mutual exclusion on modern processors would be **impossible**.

Synchronization

How are locks, semaphores, cond var, etc **implemented?**



Disabling interrupts

- Idea: lock implementation code executed in kernel mode.
 - No need for “spinlock” when we have a scheduler!
- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable
- Thread that can't take a lock goes to sleep
- Danger! What if preemption timer fires during lock implementation?
- (Again: this idea is for single-processor systems.)

Disabling interrupts

```
int value = FREE;
```

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    }  
    else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Unlock() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue;  
        place on ready queue;  
    }  
    else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

Disabling interrupts discussion

Why do we need to disable interrupts?

- Avoid preemption between checking and setting lock value
- Otherwise, two threads could think that they both have lock

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    }  
    else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```



Critical section is short
in kernel mode

What about enabling interrupts when going to sleep?

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    }  
    else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

Enable interrupts here?



☹ Release can check the queue and not wake up thread

What about enabling interrupts when going to sleep?

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    }  
    else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

Enable interrupts here?

- Release puts the thread on the ready queue, but the thread still thinks it needs to go to sleep
- Misses wakeup, but still holds lock (deadlock!) 😞😞

What about enabling interrupts when going to sleep?

```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    }  
    else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

Want to enable interrupts after sleep()

- But how?

What about enabling interrupts when going to sleep?

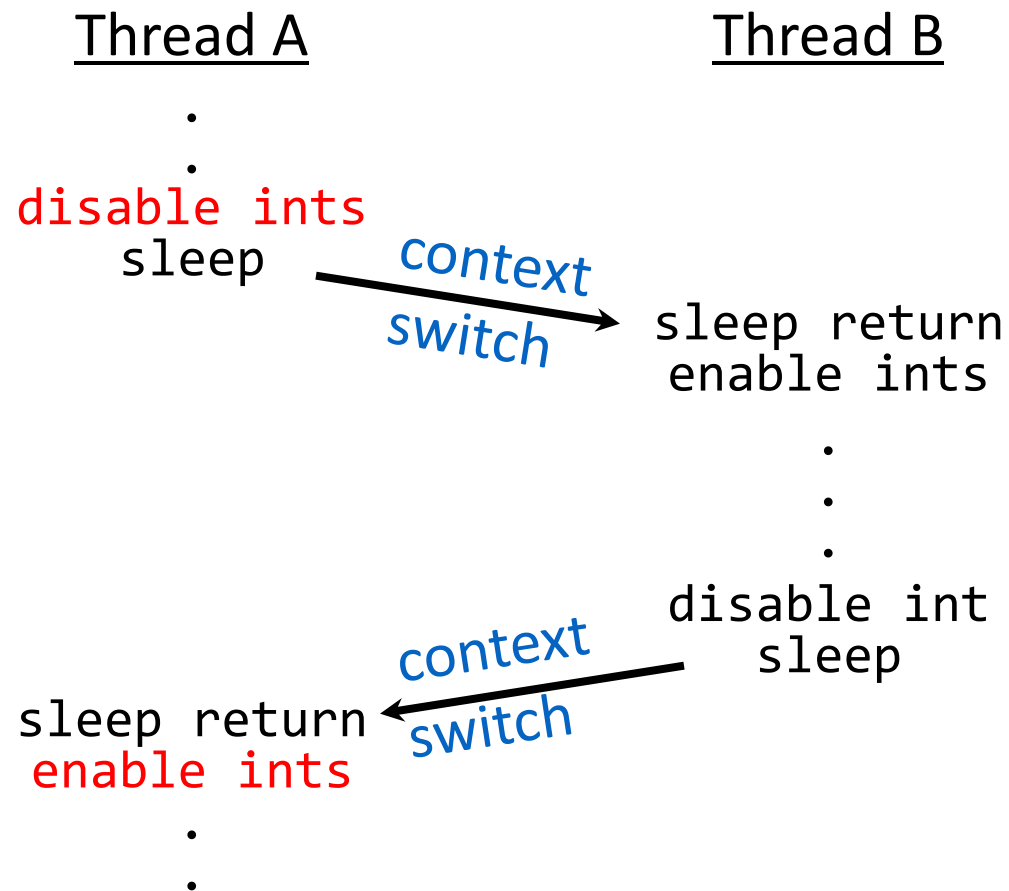
```
Lock() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    }  
    else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

Want to enable interrupts after sleep()

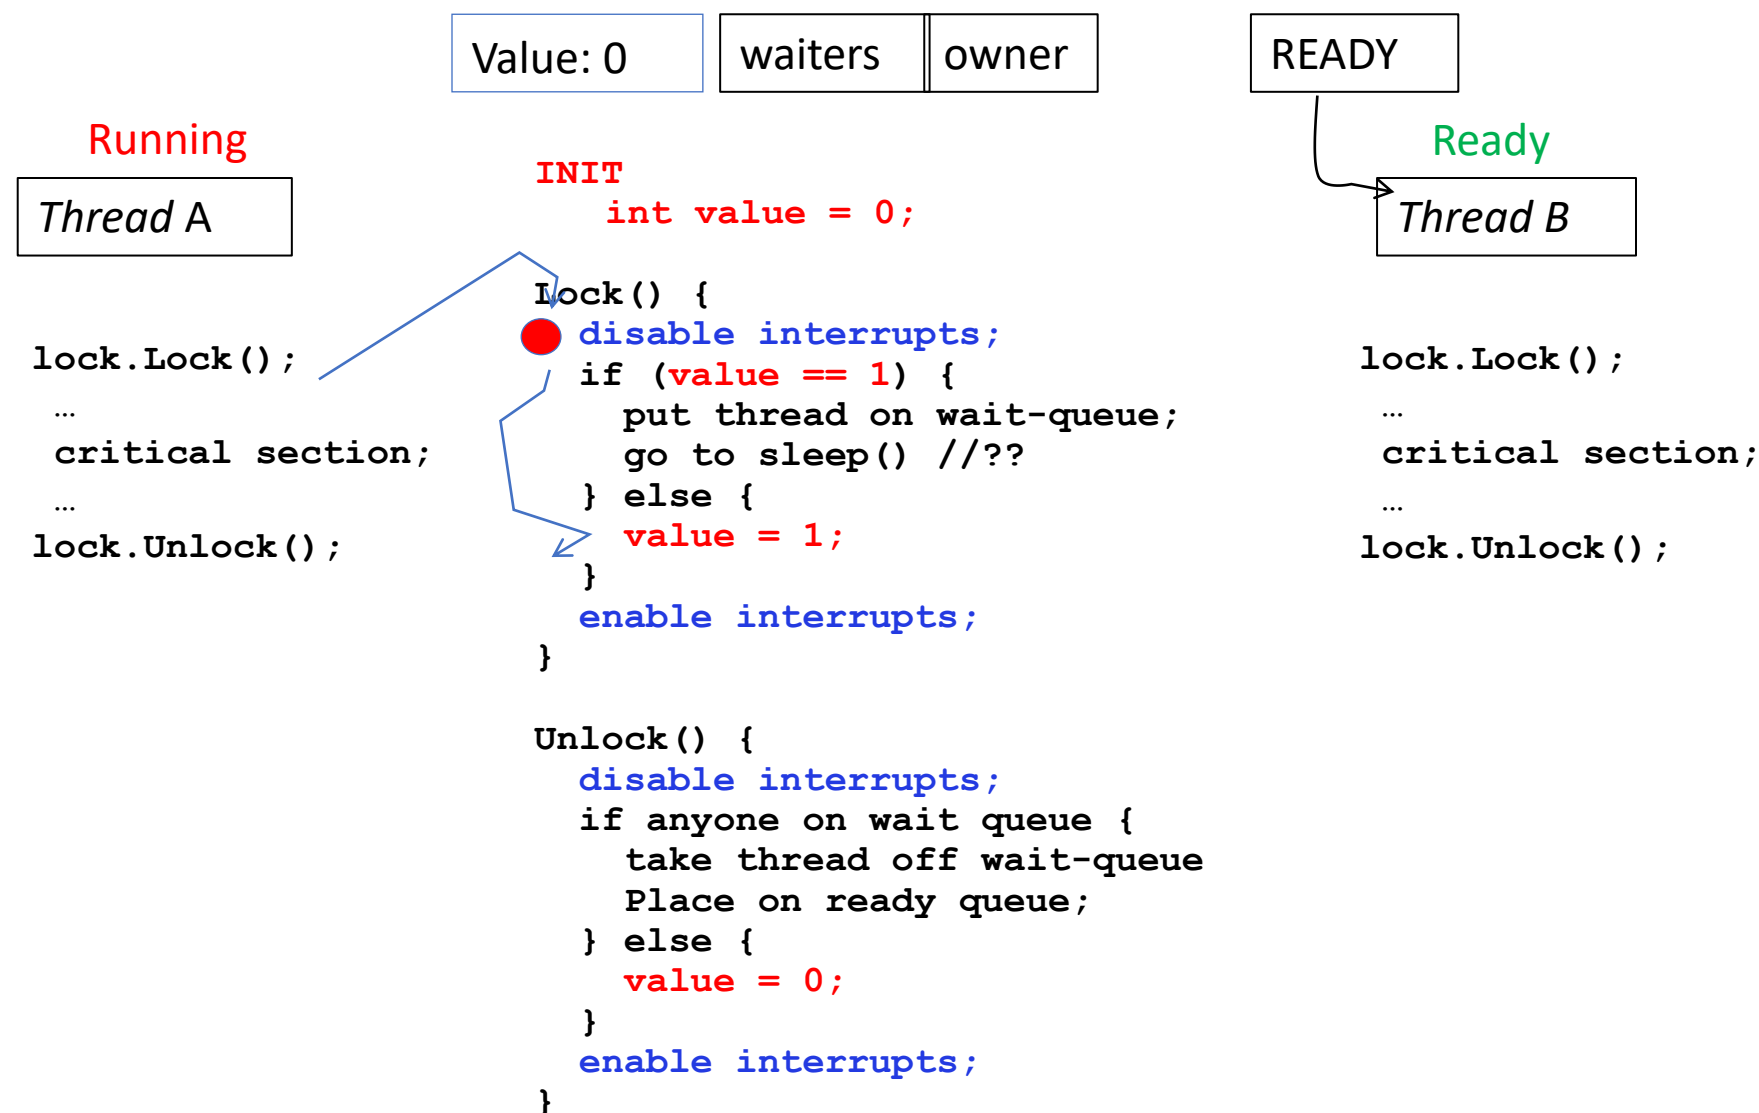
- But how? In scheduler.

In scheduler, since interrupts are disabled when you call sleep():

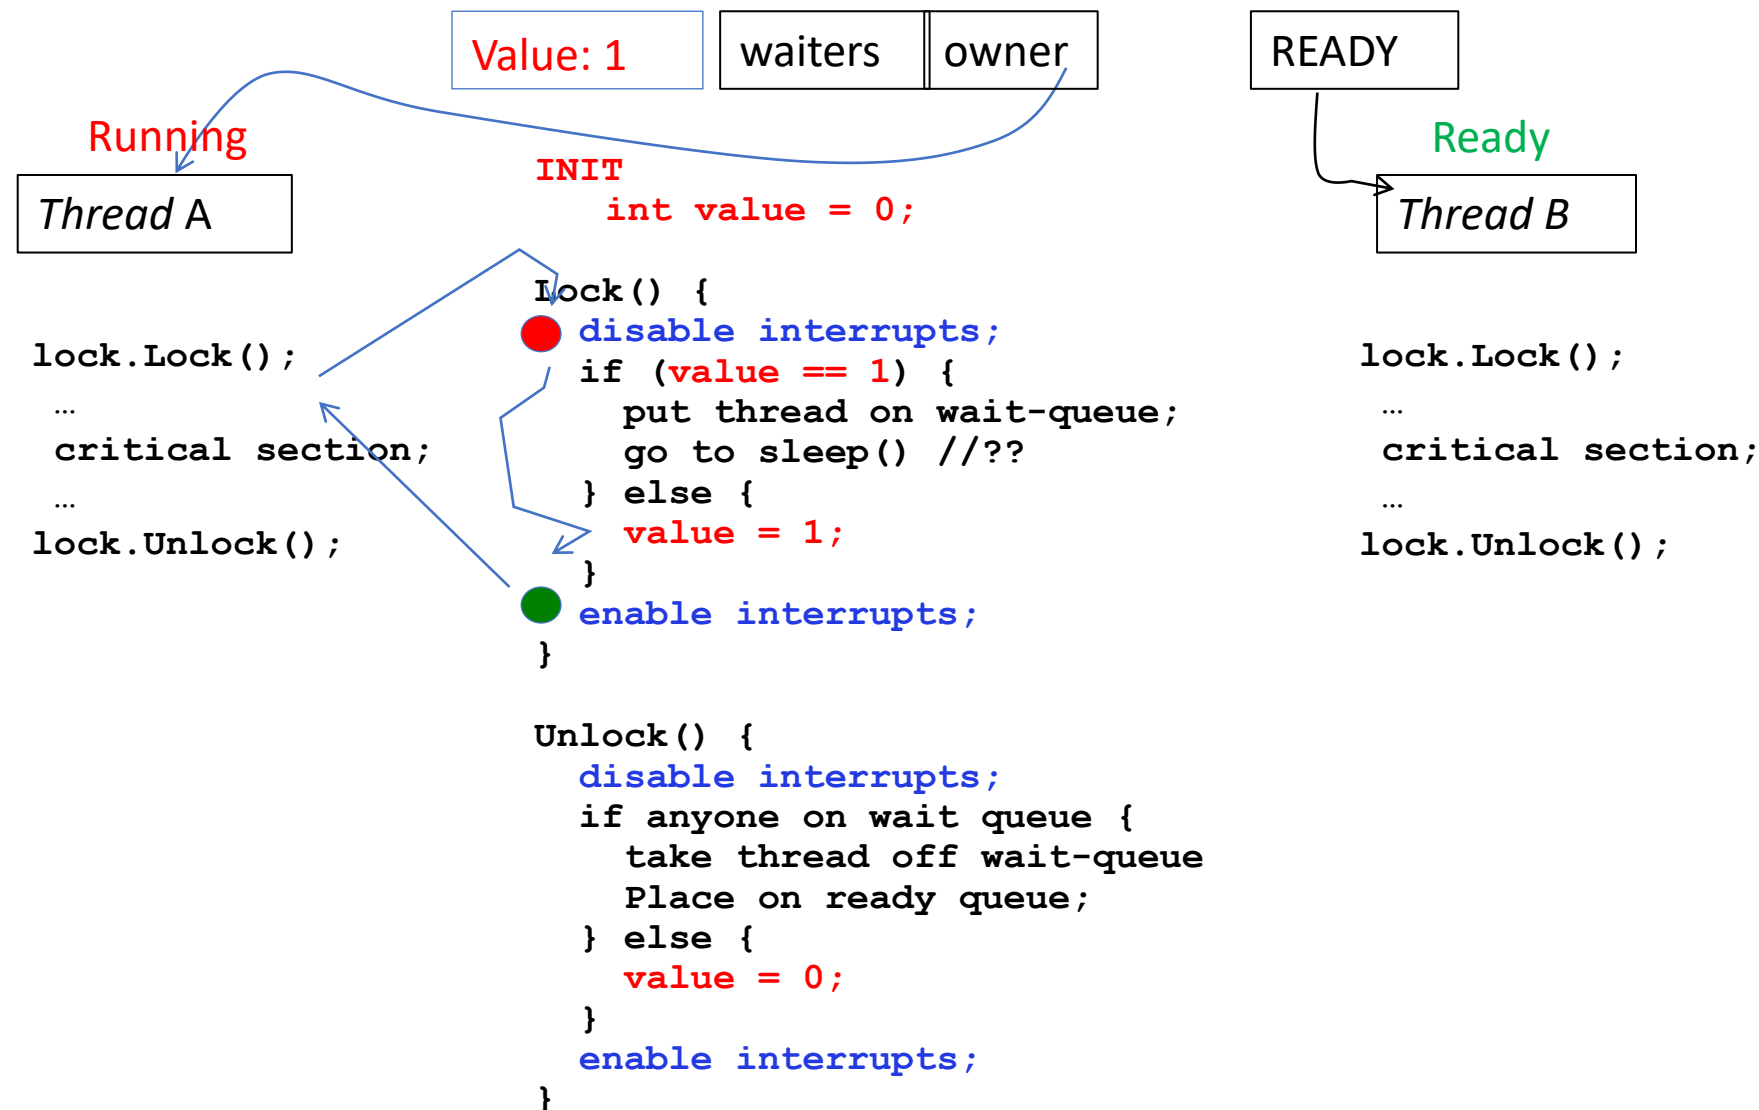
- Responsibility of the next thread to re-enable interrupts
- When the sleeping thread wakes up, returns to lock() and re-enables interrupts



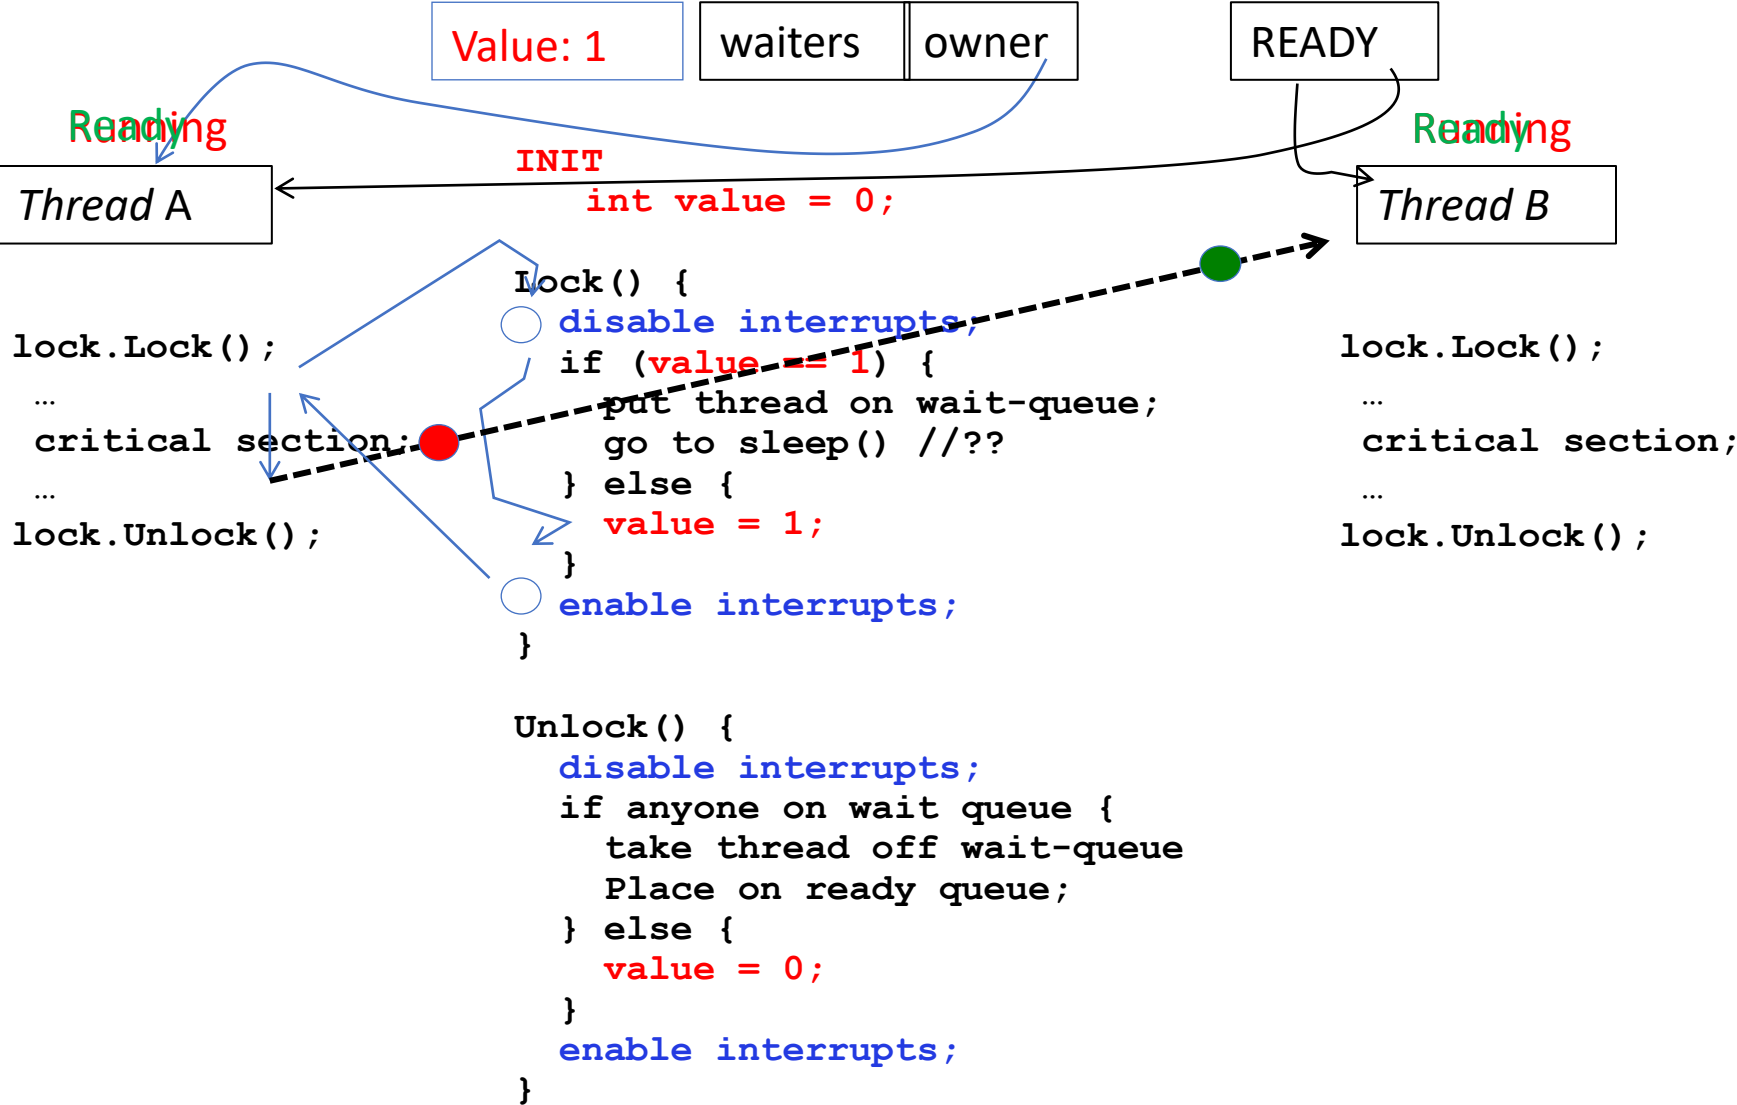
Simulation of locks with disabled interrupts



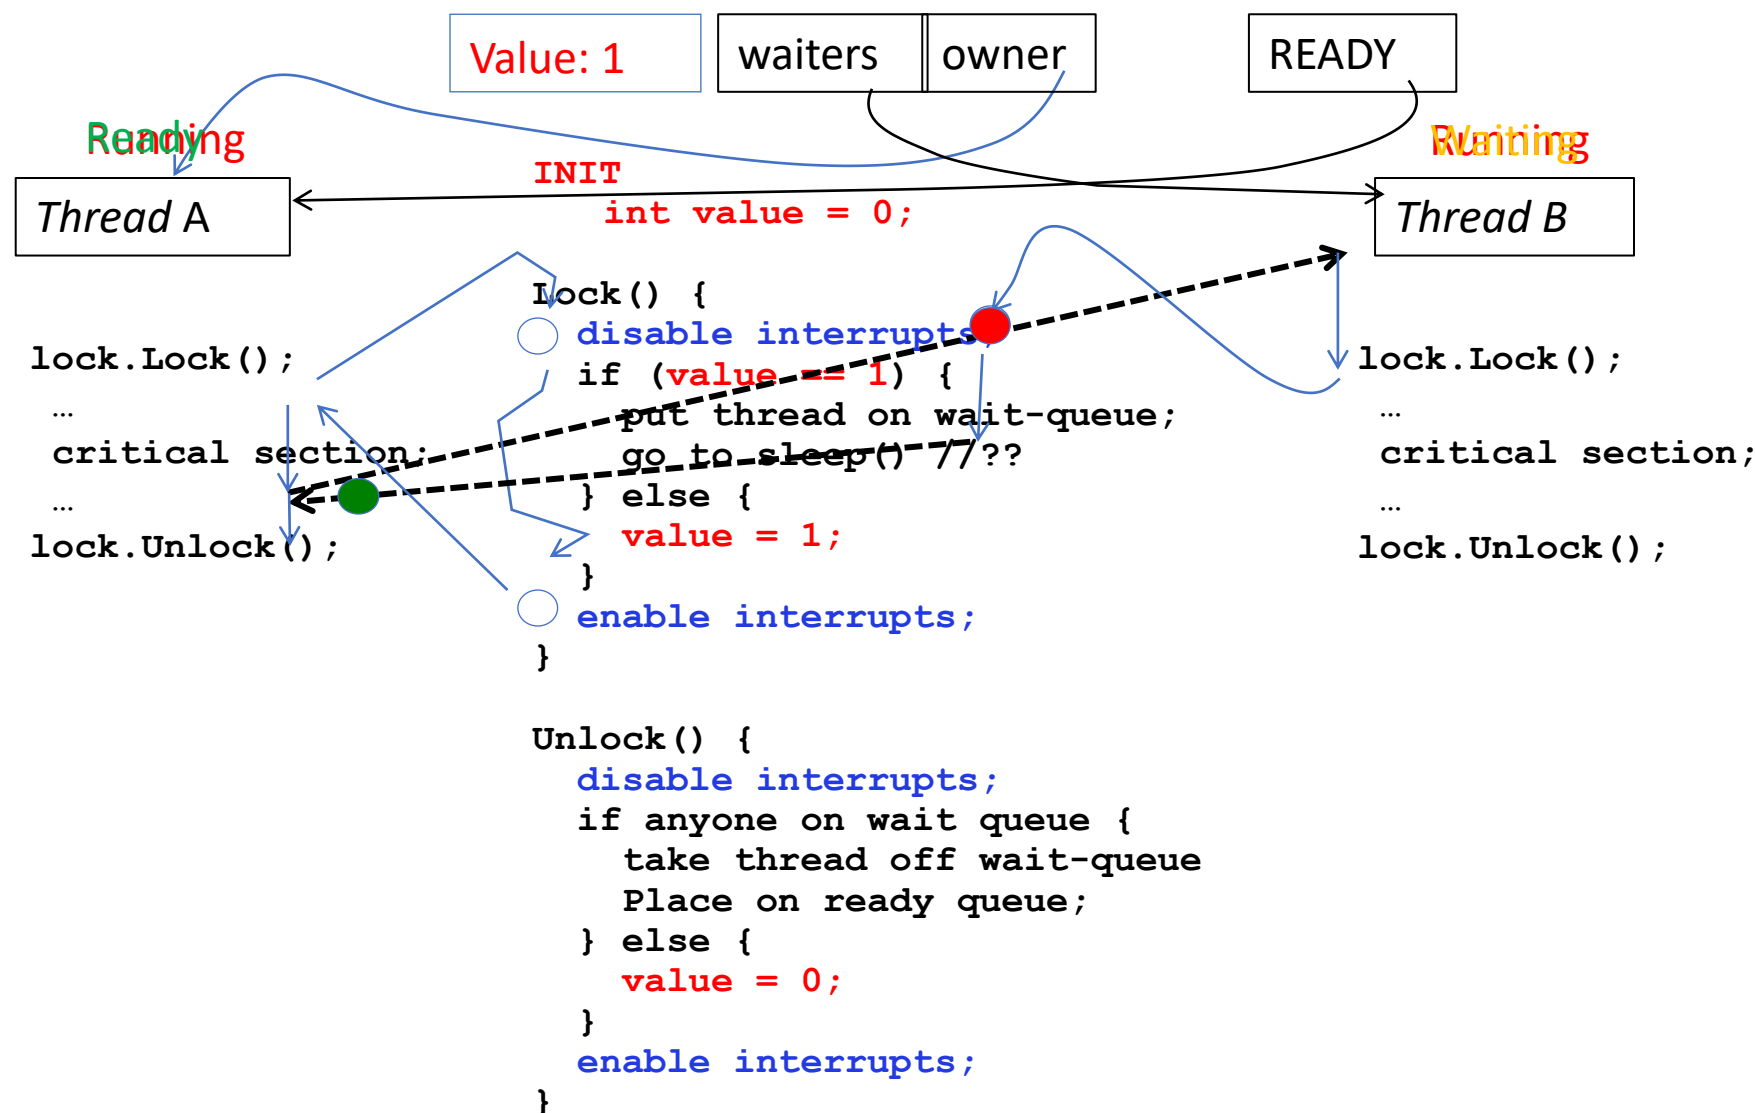
Simulation of locks with disabled interrupts



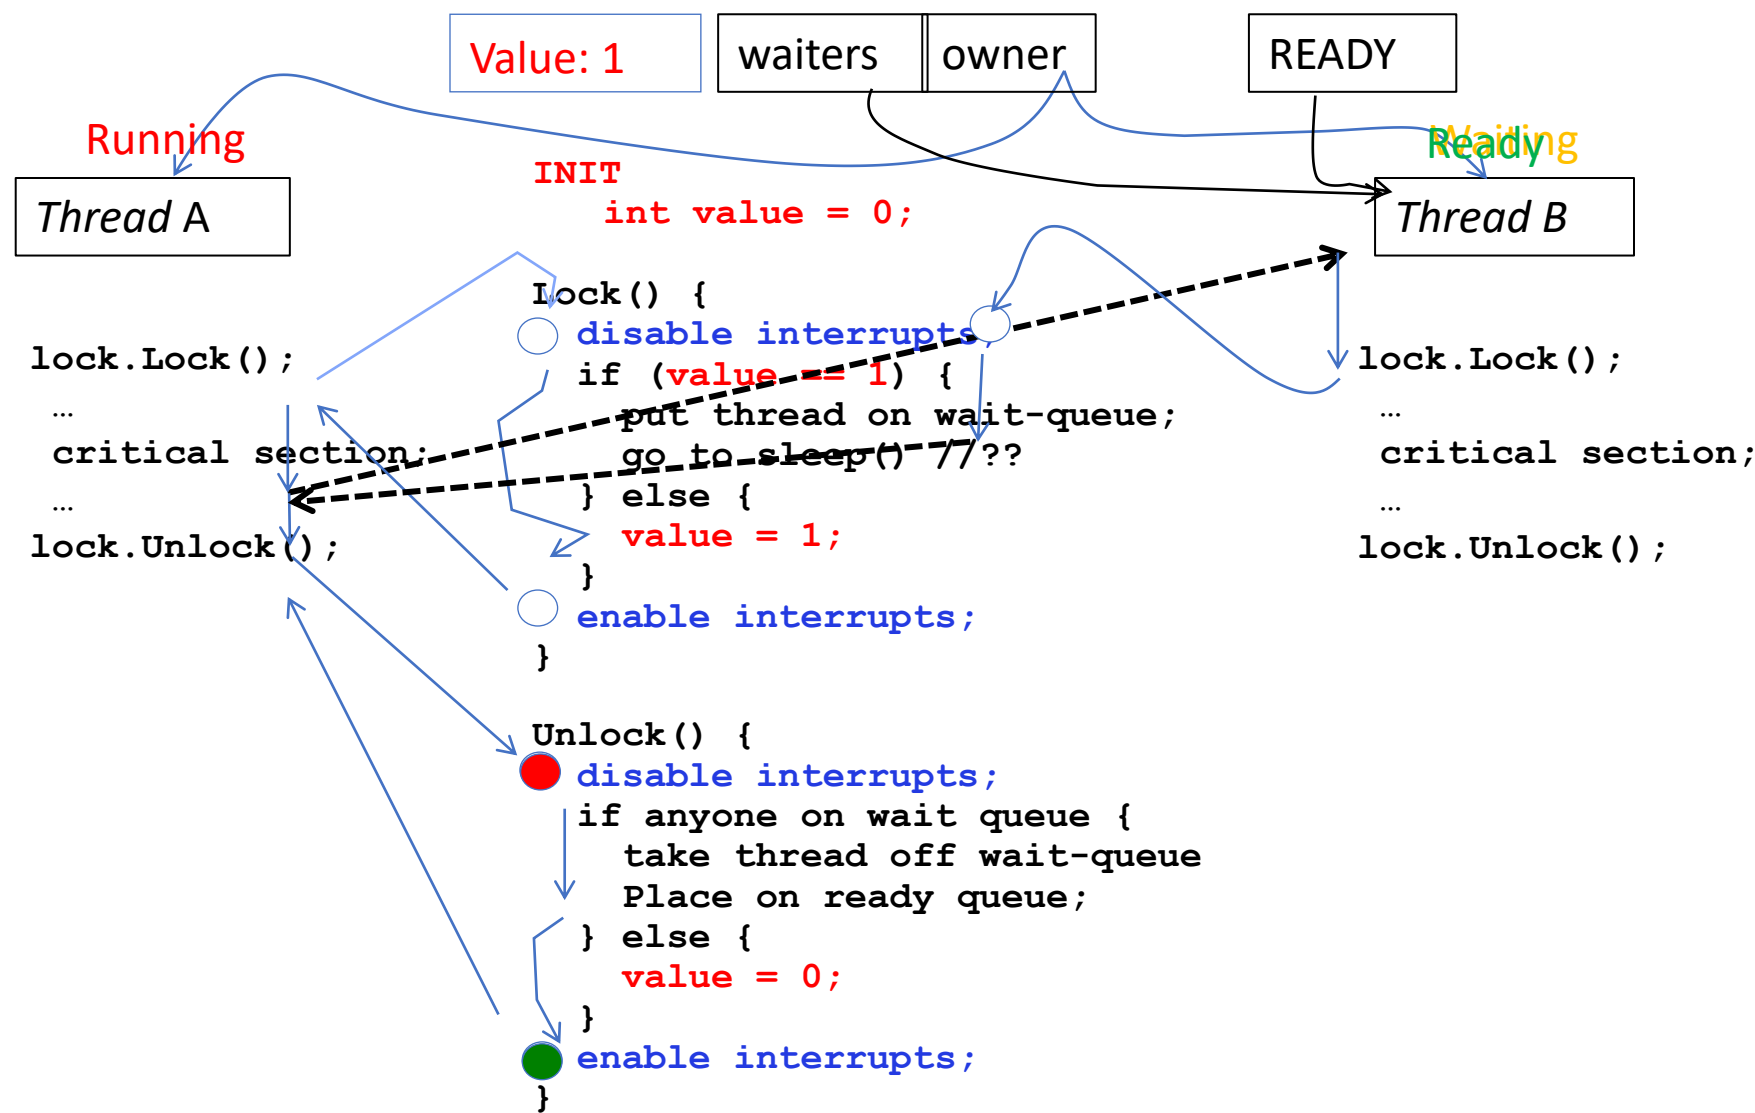
Simulation of locks with disabled interrupts



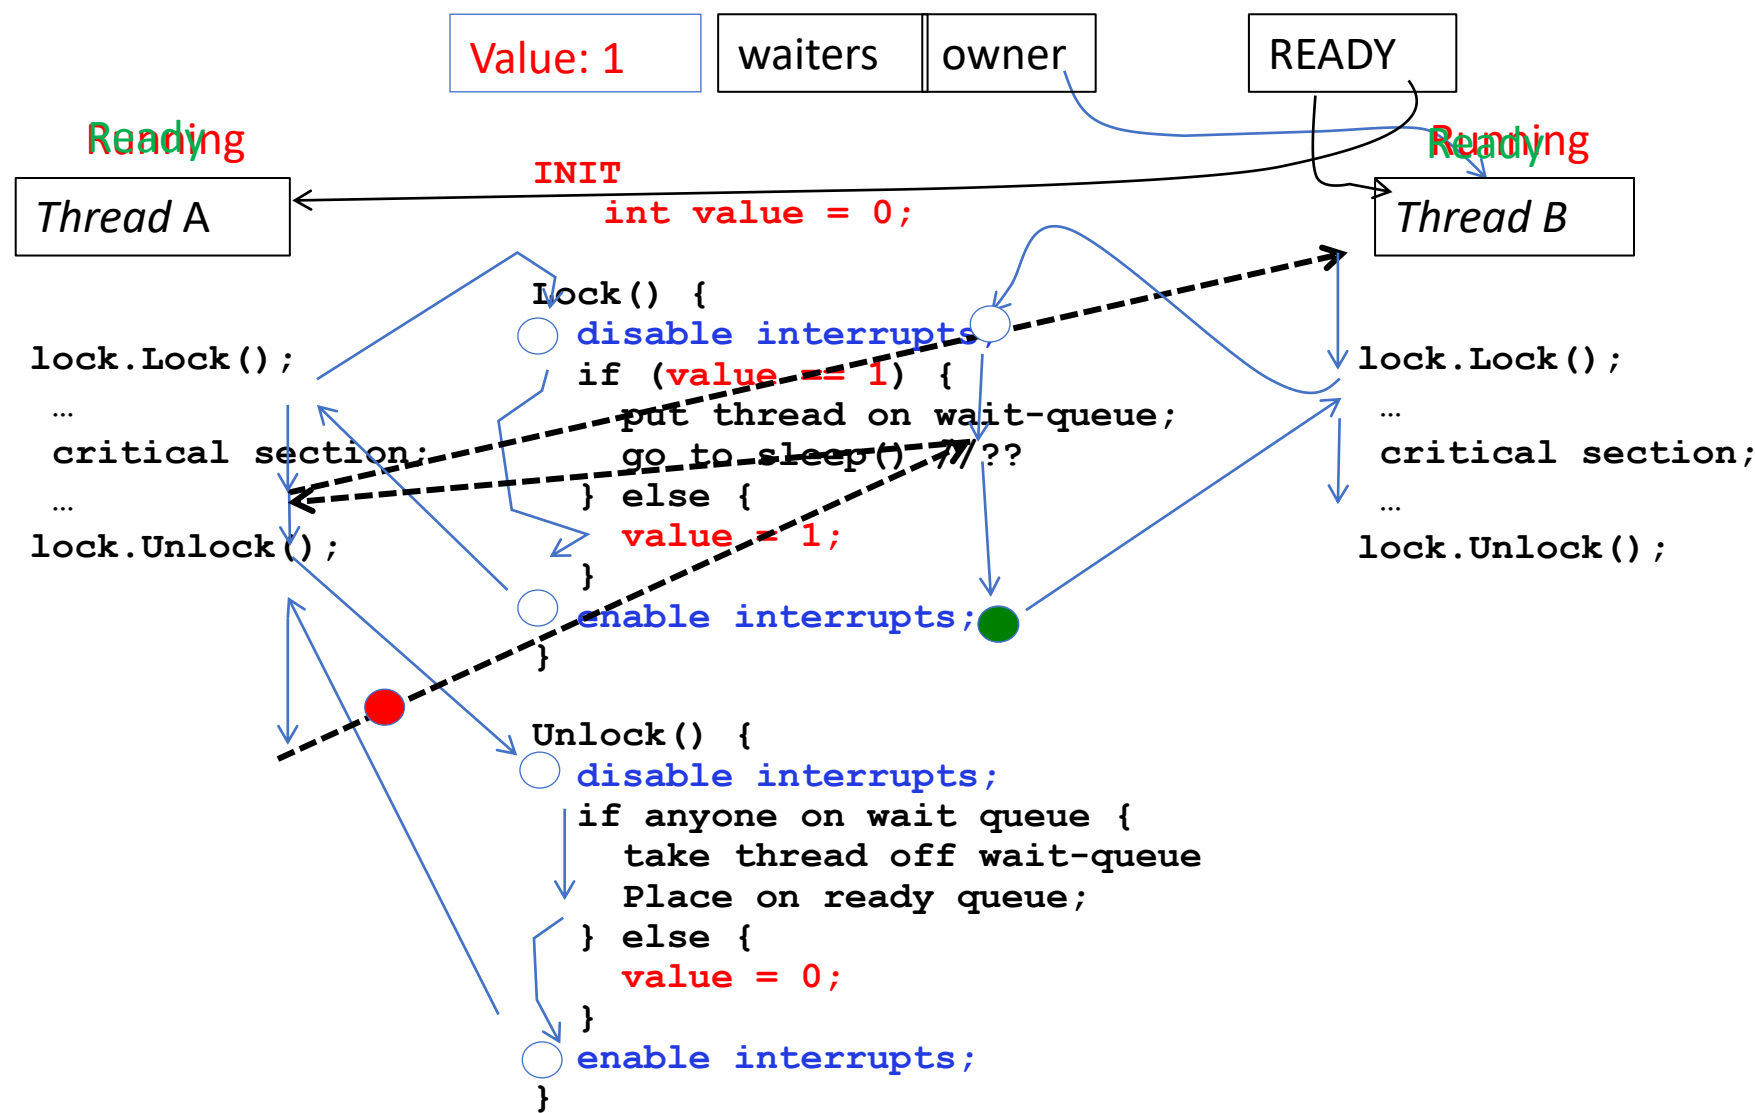
Simulation of locks with disabled interrupts



Simulation of locks with disabled interrupts



Simulation of locks with disabled interrupts



Further Optional Reading

Operating Systems: Three Easy Pieces by R. & A. Arpaci-Dusseau

Chapters 25 – 31 (inclusive)

<https://pages.cs.wisc.edu/~remzi/OSTEP/>

Reading on concurrency: Herlihy & Shavit: The Art of Multiprocessor Programming, 2nd edition.

Credits:

Some slides adapted from the OS courses of Profs. Remzi and Andrea Arpaci-Dusseau (University of Wisconsin-Madison), Prof. Willy Zwaenepoel (University of Sydney), and Prof. Maurice Herlihy (Brown University), Prof. Natacha Crooks (UC Berkeley).

