

Graded Exercises Solutions

COMP 310 / ECSE 427 Winter 2025

1 Syscalls [10 points]

1.1 Syscall Wrappers [5 points]

When a process makes a system call through the kernel API library with a pointer as an argument, does the validity of that pointer need to be checked in the library, in the kernel, or in both places? Justify your answer.

There are three possible answers: both, just library, or just kernel. “Both” or “just kernel” are both acceptable answers with proper justification.

The validity of a pointer must be checked before it can be dereferenced. It is especially important in the kernel to check the validity of a userspace pointer before dereferencing it in the kernel, as we have to do a virtual address translation and a failure to check validity could give the user access to protected memory!

If the library wrapper has shortcut paths, where (perhaps via caching or some other mechanism) it can sometimes produce results without a system call, then any pointers that it needs to dereference to produce or return that result need to be checked for NULL or other problems, or else there is a risk of segfaults. (Note that we cannot check for virtual address validity in the library, since it also runs in userspace.)

1.2 Forking [5 points]

How many times will the below program print `hello5`? Explain why (e.g. by drawing the process tree).

```
1  int main() {
2      if (fork() != 0) {
3          if (fork() != 0) {
4              printf("hello1\n");
5          } else {
6              fork();
7          }
8          printf("hello2\n");
9      } else {
10         if (fork() != 0) {
11             printf("hello3\n");
12         } else {
13             fork();
14         }
15         printf("hello4\n");
16     }
17     printf("hello5\n"); // <-- this one!
18 }
```

Each call to `fork()` is reached by exactly one process. Exactly two processes (parent and child) return from each call. Therefore each call generates one process. We begin with one process, there are five calls, so we end with 6 processes. All 6 processes reach the final line of code, so `hello5` is printed 6 times.

2 Synchronization [10 points]

Assume that a finite number of resources of a single type must be managed in a multi-threaded program. Threads acquire a number of these resources and – once finished – return them. The following program segment is used to manage a finite number of instances (`MAX_RESOURCES`) of an available resource. When a thread wishes to obtain a number of resources, it invokes the `decrease_count` function. When a thread wants to return a number of resources it calls the `increase_count` function.

```
1 #define MAX_RESOURCES 5
2 int available_resources = MAX_RESOURCES;
3
4 int decrease_count(int count) {
5     if (available_resources < count) {
6         return -1;
7     } else {
8         available_resources -= count;
9         return 0;
10    }
11 }
12
13 int increase_count(int count) {
14     available_resources += count;
15     return 0;
16 }
```

2.1 A Race Everyone Loses [1 point]

Unfortunately, the above program fragment causes race conditions. Identify the data involved in the race condition(s).

`available_resources`. The `count` arguments are not involved in any races because they are local to the functions and therefore also thread-local.

2.2 Don't Bet on the Trifecta [2 points]

Clearly identify the location(s) in the code where the race condition(s) occur.

The references to `available_resources` on lines 5, 8, and 14. All three references are involved in races, because read-write races between line 5 and 8 (for example) can still cause races.

In particular, the danger of this race is that we may perform a check on line 5 and find that there are not enough resources at the same time as another thread increases the count, in which case the `decrease_count` function will fail when it could have succeeded. This is the bug that needs to be fixed in the next part.

2.3 Save the Track [7 points]

Fix the race condition. You can either write **pseudo-code** to replace the program above, or clearly explain how to modify the code. (You can refer to C library functions and values seen in class.)

We need to protect access to the shared resource with mutual exclusion in order to fix the bug described above. In particular, we need a mutex used by both functions. It needs to be locked between lines 13/14 and unlocked between 14/15. To protect the `decrease_count` function, it needs to be locked before we perform the check, so between lines 4/5, and unlocked before the function returns. It must be unlocked. So we need to unlock between lines 5/6 and also between 8/9. Alternatively, we can set a “return value” variable and move the return statements out of the conditionals so that we only have to unlock the mutex in one place.

3 Compaction [5 points]

This question asks you to do a bit of math. Specifically, we want to see that you understand *what compaction is doing* and how that affects the scenario.¹ We do not need to see that you can accurately do arithmetic. Use a calculator if you wish, and give an approximate answer (say, two digits of precision).

3.1 Turbocharge the matter compressor [5 points]

(Yes, that is a Futurama joke.)

A memory management system eliminates holes by compaction. Assume that a random distribution of holes and many segments means that very nearly the entire memory must be copied. Assume we can read an 8-byte value from physical storage in 4 nanoseconds (ns) and that we can also write 8-byte values in 4ns.

How long does it take to compact 4GB of memory? Explain your reasoning.

Hint: does a copy require only a read, only a write, or both?

The assumption says we need to copy all 4GB of memory. A copy requires both a read and a write. So it takes $4 + 4 = 8\text{ns}$ to copy 8 bytes.

Total time required is therefore $4\text{GB} / 8\text{B} * 8\text{ns} = 4\text{Gns}$ (aka 4 seconds).

4 Virtual Memory: Segmentation [10 points]

For this question, all numbers prefixed by 0b and in that font are in binary. Therefore, 0b100 is the decimal number 4.

A CPU boasts 5-bit virtual addresses and 16 bytes of byte-addressable physical memory.² There are three segments currently mapped, with the following segment table:

Seg#	Base	Bound	R W
0	0b1100	0b100	0 1
1	0b0000	0b000	0 0
2	0b0001	0b011	1 1
3	0b1000	0b010	1 0

4.1 Cartography [1 point each]

Give either the corresponding physical address, or the phrase “segmentation fault,” for each of the following accesses. If the access faults, briefly explain why. (You may draw a diagram of the physical memory if that makes it easier to explain.)

- 0b00000 read segment 0, no permission. segmentation fault
- 0b11001 read segment 3, ok. 0b1001
- 0b01000 read segment 1, either reason. segmentation fault
- 0b10010 read segment 2, ok. 0b0011
- 0b11010 read segment 3, out of bounds. segmentation fault
- 0b00111 write segment 0, out of bounds. segmentation fault
- 0b10001 write segment 2, ok. 0b0010
- 0b00000 write segment 0, ok. 0b1100
- 0b11000 write segment 3, no permission. segmentation fault
- 0b10100 write segment 2, out of bounds. segmentation fault

¹Note also that these numbers are not realistic.

²Quite the boast, huh? Don't tell them.

5 Virtual Memory: Paging [15 points]

We will consider two machines in this question. (There is a third part on the next page.) The first machine has 32-bit virtual addresses. The machine supports paging with a 4KB page size and a 2-level paging scheme. The page directory³ contains 1024 entries.

5.1 This Sounds Oddly Realistic [2 points]

How many entries are there in a single inner-level page table? Justify your answer.

How many bits for each fragment of the virtual address? Well, the page directory has 2^{10} entries, so the page directory index is 10 bits. There are $4\text{KB} = 2^{12}$ per page, so the offset is 12 bits. That leaves 10 bits for the page table index. Therefore there are 1024 entries in an inner-level page table.

Unlike the previous question, these numbers are realistic. In fact, they are so realistic, that these are actually the numbers used in most 32-bit processors with support for paging. That includes old x86 processors (but not modern AMD64 processors, which are 64-bit versions of x86), as well as 32-bit ARM processors.

Knowing that the numbers are realistic, we can arrive at the right answer in a few other ways as well: for example, knowing that page tables and page directories should be the same size in a real system. Alternatively, we could recall from lecture that we **really** want the page tables/directories to fit exactly into a frame. Our 32-bit machine should have 4 byte entries. A 4KB frame can hold $4\text{KB} / 4 = 1\text{K}$ entries of that size – 1024 is the right answer.

An exam question will not require such reasoning, just like this question does not require it, but it may definitely help.

5.2 Spatial Locality [5 points]

A program running on this first machine is only sparsely using its address space. The addresses between 0 (inclusive) and 18MB (exclusive) are mapped. Additionally, the addresses between 90MB and 117 MB are also mapped.

What is the minimum number of inner-level page tables needed? Justify your answer and take special care to consider whether or not 90 and 117 are on page directory boundaries. (Again, we want to see that you know how to approach the problem, not that your arithmetic is accurate.)

Each page table has 1K entries that each handle a 4K page. That means each page table handles $1\text{K} \times 4\text{K} = 4\text{M}$ virtual addresses.

We need to cover address 0-18M at the bottom of the address space, which will require 5 tables; the ones for 0-4M, 4-8M, 8-12M, 12-16M, and 16-20M. We won't use all of the last page table, but we need *some* of it, so it must be allocated.

At the top of the range we need to cover $117\text{M} - 90\text{M} = 27\text{M}$ addresses. Unfortunately it's not as simple as $28\text{M} / 4\text{M} = 7$ tables. It's not that simple, because there is no page table that covers 90M-94M. 90M is not divisible by 4M, but the page table boundaries are each 4M as computed above! The first page table needed for this range will be 88M-92M. The last page table will be 116M-120M. Count them – there are 8.

$5 + 8 = 13$ page tables.

³Recall “page directory” is the name for the outer-level page table.

5.3 Chasing Stars [8 points, 2 each]

This question uses some hexadecimal numbers, formatted like `0xf3`.

A second machine has 6-bit addresses (both virtual and physical). It has 64 bytes of byte-addressable physical memory. The designers heard about the cool two-level paging technique and decided to implement it. The page size is 4 bytes; each entry is a single byte. Like usual, the page tables and the page directory are themselves page-sized. Therefore, each one contains 4 entries. Overall, this was probably a bad idea, because the page tables take up a large proportion of the machine's small memory. The physical memory contains only $64/4 = 16$ frames. Therefore, page table entries use only the the bottom four bits (one hexit) to store a frame number. The other bits are used for metadata.⁴ **If the metadata bits are all zero, the entry is invalid.**

At right, the first 8 frames of physical memory are shown. The page directory occupies **frame number 2**. (Frame numbers are 0-indexed, so the frame numbered 2 is **not** the second frame!) For each of the following **virtual** addresses, perform the page table walk to determine either the corresponding physical address or “address invalid.” Briefly justify each step of the walk by giving the physical address of **each** lookup. Giving only the final translated address is not enough for full credit.

You may give addresses either in binary or as (hexadecimal) “frame number, offset” pairs. For example, `0b110110` and `d,2` denote the same address.

Hint 1: The virtual addresses are given in binary so that you can easily separate them into parts. There are two bits for each of the offset, page table index, and page directory index. (You could have deduced this from the other given information.)

Hint 2: The top 4 bits of all of the physical addresses in a frame are the same. That is not a coincidence! That's the frame number in binary.

1. `0b000000`
2. `0b110010`
3. `0b101101`
4. `0b010111`

By the way, there is a secret hidden somewhere in the virtual address space. See if you can find it!

1. `0b001000` (or `2,0`) invalid PDE
2. `0b001011` (or `2,3`) → `0b001100` (or `3,0`) invalid PTE
3. valid, `0b001010` (or `2,2`) → `0b010111` (or `5,3`) → `0b110101` (or `d,1`)
4. valid, `0b001001` (or `2,1`) → `0b000001` (or `0,1`) → `0b011111` (or `7,3`)

Note that the final address in the last case contains the value 0. This is a data value in the user's address space that happens to be 0. Nothing is invalid there.

Food for thought: what would happen if one of the page directory entries pointed back to the page directory? Might this be useful for something?

Addr	Value
<code>0b000000</code>	<code>0xd1</code>
<code>0b000001</code>	<code>0xf7</code>
<code>0b000010</code>	<code>0x6f</code>
<code>0b000011</code>	<code>0x00</code>
<code>0b000100</code>	<code>0x53</code>
<code>0b000101</code>	<code>0x65</code>
<code>0b000110</code>	<code>0x63</code>
<code>0b000111</code>	<code>0x72</code>
<code>0b001000</code>	<code>0x00</code>
<code>0b001001</code>	<code>0xf0</code>
<code>0b001010</code>	<code>0xf5</code>
<code>0b001011</code>	<code>0xf3</code>
<code>0b001100</code>	<code>0x00</code>
<code>0b001101</code>	<code>0xca</code>
<code>0b001110</code>	<code>0x44</code>
<code>0b001111</code>	<code>0xc8</code>
<code>0b010000</code>	<code>0xaf</code>
<code>0b010001</code>	<code>0x05</code>
<code>0b010010</code>	<code>0xf5</code>
<code>0b010011</code>	<code>0x45</code>
<code>0b010100</code>	<code>0xd9</code>
<code>0b010101</code>	<code>0x00</code>
<code>0b010110</code>	<code>0xfc</code>
<code>0b010111</code>	<code>0x5d</code>
<code>0b011000</code>	<code>0xef</code>
<code>0b011001</code>	<code>0x59</code>
<code>0b011010</code>	<code>0xee</code>
<code>0b011011</code>	<code>0xd6</code>
<code>0b011100</code>	<code>0x65</code>
<code>0b011101</code>	<code>0x74</code>
<code>0b011110</code>	<code>0x21</code>
<code>0b011111</code>	<code>0x00</code>
...	...

⁴For example, read/write permissions. But we will ignore that here.